Inhibitory Attention for Clinical Long-Context Reasoning: Characterizing and Mitigating Lost-in-the-Middle Effects in EHR Processing

Anonymous Author(s)
Affiliation
Address
email

Abstract

Electronic health records now routinely exceed 100,000 tokens per patient, placing them at the frontier of long-context language model capacity. Yet large language models exhibit a systematic positional retrieval bias—the lost-in-the-middle (LitM) effect—in which information located near the center of a long input context is retrieved far less reliably than information near the edges. For clinical applications, this bias is not a minor benchmark quirk: a discontinued medication recorded two years ago, a critical lab value from the middle of a patient's record, or an old diagnosis that contradicts a current treatment plan may represent the most clinically consequential fact in the entire note. We term this the *clinical lost-in-the-middle* (CLitM) problem, provide its first systematic empirical characterization across clinical fact domains using MedAlign, and compare context selection strategies as remedies. Across 2,196 instruction-response pairs and six language models, we observe a 21.9 percentage-point gap between peak accuracy (59.5% [95% clustered CI: 46.3, 71.0] at the 20–30% position decile) and trough accuracy (37.6% [23.2, 52.5] at 70–80%; gap clustered CI: [17.7, 45.2] pp); 67.8% of reference answers reside between the 10th and 90th percentiles of the EHR timeline, placing them within the CLitM accuracy trough. We introduce Query-Conditioned Clinical Suppression (QCCS), a lightweight character n-gram query-conditioned context selection gate, and evaluate it against BM25, BM25 with section-header filtering, dense retrieval (sentence-transformers), and cross-encoder reranking in a multi-arm proof-of-concept experiment (N=83 held-out instructions). End-to-end instruction-following experiments with Qwen2.5-7B-Instruct (16k-token context window) show that QCCS substantially outperforms all five comparators, evaluated by LLM-as-judge semantic scoring (primary: Claude Haiku; second judge: claude-sonnet-4-6; Cohen's κ=0.767 for QCCS, the arm with meaningful positive-class variation): QCCS 16.7% [95% CI: 3, 30] vs. BM25 3.3% [0, 10] vs. cross-encoder 0.0% [0, 0] vs. dense 0.0% [0, 0] vs. full context 6.7% [0, 17] for middle-position instructions; overall, QCCS reaches 25.3% [17, 35] vs. $\leq$3.6% for all retrieval-only comparators. This end-to-end advantage is not attributable to higher retrieval recall: at k=20, BM25 retrieves the gold-standard evidence sentence in 98.8% of test instructions (QCCS: 34.9%); cross-encoder reranking achieves 96.4% overall; dense retrieval 94.0%. A conditional accuracy analysis reveals that Stage-2 accuracy remains $\leq$2.6% for all retrieval-based comparators even in the 74–82 instructions where those arms successfully retrieve the gold-evidence sentence (by lexical overlap), while QCCS achieves 25.0% accuracy in the 52 of 83 instructions where it does *not* retrieve the gold sentence—establishing that **in this proof-of-concept evaluation,**

Submitted to 40th Conference on Neural Information Processing Systems (NeurIPS 2026). Do not distribute.

query-aligned context selection is a stronger predictor of EHR instruction-following accuracy than gold-sentence retrieval recall. A no-query ablation confirms that query conditioning accounts for +23.8 percentage points of middle-position recall. Code is available at `[anonymized for review]`.

1 Introduction

Modern electronic health records are among the longest structured documents routinely encountered by language models. A single longitudinal patient record from an integrated delivery network can exceed 100,000 tokens when all encounter notes, laboratory results, medication histories, and social history entries are concatenated in chronological order. Recent advances in transformer architectures—extended context windows, flash attention, ring attention—have made it technically possible to feed an entire patient history into a single forward pass. The critical question is whether the model can faithfully retrieve information from any position in that history, or whether its attention mechanism systematically privileges certain positions at the expense of others.

The lost-in-the-middle (LitM) phenomenon, first characterized by Liu et al. [2024] in multi-document question answering, establishes that language models are significantly less accurate at retrieving information from the center of a long input context than from the edges. The resulting U-shaped accuracy curve reflects a fundamental property of standard softmax attention: the mechanism attends preferentially to tokens near the beginning and end of the sequence, leaving a systematic gap in the middle. Liu et al. [2024] demonstrated gaps as large as 20 percentage points across a range of state-of-the-art models on non-clinical benchmarks. Whether an analogous gap exists in clinical EHR contexts—and whether the clinical stakes of this gap warrant architectural rather than prompt-engineering solutions—has not been studied.

This gap has particular clinical urgency. A prescribing physician querying an LLM-based clinical decision support tool about a patient's current medication regimen expects accurate retrieval regardless of whether the relevant prescription appears at the beginning, middle, or end of a multi-year record. Existing mitigation strategies—retrieval-augmented generation (RAG), prompt compression via LLMLingua Pan et al. [2024], and KV cache eviction methods such as H_2O Zhang et al. [2023] and SnapKV Li et al. [2024]—address related problems but not the core architectural failure. Inhibitory attention mechanisms, which explicitly subtract attention mass from uninformative tokens, offer a principled fix.

This paper makes four contributions:

1. **Clinical LitM characterization.** We provide the first systematic empirical characterization of the CLitM effect in real clinical EHR contexts, using needle-in-a-haystack experiments over MedAlign Fleming et al. [2024] across four clinical domains (medications, diagnoses, laboratories, and social determinants of health) and six language models, establishing a 21.9 percentage-point peak-to-trough accuracy gap.
2. **Inhibitory attention on structured EHR prediction.** We evaluate Differential Transformer Ye et al. [2025] (DiffAttn) and the full QCCS-gated integration (QCCS-DiffAttn, Eq. 3) versus standard attention and KV compression baselines on four EHRSHOT Wornow et al. [2023] laboratory prediction tasks, finding consistent gains on class-imbalanced tasks with DiffAttn (+6.1 pp AUROC on anemia; +4.2 pp on hyperkalemia), and finding that per-token QCCS gating (QCCS-DiffAttn; Eq. 3) does not improve over scalar inhibition: gate training collapses to chance on hyperkalemia (AUROC 0.500) and AUROC falls below scalar DiffAttn on all four tasks.
3. **Query-Conditioned Clinical Suppression (QCCS).** We introduce QCCS, a lightweight character n-gram query-conditioned context selection gate that scores EHR sentences for relevance to the clinical query. The gate is trained on MedAlign instruction-response pairs using lexical overlap labels. We also evaluate (Eq. 3) the integration of the gate score into a Differential Transformer as a per-token inhibition weight (QCCS-DiffAttn).
4. **Retrieval baselines and end-to-end evaluation.** We compare QCCS against BM25, BM25 with section-header filtering, dense retrieval, and cross-encoder reranking on a clean held-out test split (74 patients, 83 unique instructions). All lexical, neural, and dense retrievers substantially outperform QCCS in gold-answer recall (BM25 98.8% overall; dense 94.0%; vs. QCCS 34.9%). LLM-as-judge end-to-end evaluation (seven arms; Qwen2.5-7B-Instruct; second independent

judge $\kappa = 0.767$ on QCCS arm) shows QCCS outperforms all comparators (25.3% [17, 35] overall vs. $\leq$3.6%; 16.7% [3, 30] middle vs. $\leq$6.7%). Conditional accuracy analysis establishes that BM25 achieves only 2.6% Stage 2 accuracy even when the gold sentence is retrieved (78/83 instructions), while QCCS achieves 25.0% *without* the gold sentence (52/83 instructions), establishing that query-conditional context alignment may matter more than retrieval recall for EHR instruction-following accuracy in this setting.

2 Related Work

Lost in the Middle. Liu et al. [2024] characterized LitM in multi-document QA and key-value retrieval, finding a consistent U-shaped accuracy curve (15–25 pp trough) across models from GPT-3.5-Turbo to Claude-2. All prior LitM characterization uses synthetic benchmarks; no study has examined the phenomenon in clinical EHR contexts.

Inhibitory attention mechanisms. McDougall et al. [2023] established that a subset of GPT-2 attention heads perform *copy suppression*—reducing logit contributions for contextually frequent tokens. The Differential Transformer Ye et al. [2025] generalizes this into a trainable architecture via differential softmax maps over separate query-key projections, scaled by a learnable parameter λ:

$$A_{\text{diff}} = \text{softmax}\left(\frac{Q_1 K_1^\top}{\sqrt{d}}\right) - \lambda \cdot \text{softmax}\left(\frac{Q_2 K_2^\top}{\sqrt{d}}\right) \tag{1}$$

Lv et al. [2024] extend this to explicit negative attention weights. Adaptively Sparse Transformers Correia et al. [2019] apply α-entmax transformations to produce hard zeros in attention maps—functionally inhibitory by suppressing zero-weight tokens—but target sparsity rather than positional bias correction. None of these architectures have been evaluated on clinical EHR data or demonstrated to mitigate CLitM in instruction-following tasks.

Context compression and query-conditioned selection. KV cache compression methods Zhang et al. [2023], Li et al. [2024], Pan et al. [2024], Bui et al. [2025], Zhou et al. [2024] improve memory efficiency but operate post-attention: middle-context tokens with low attention due to positional bias are evicted, exacerbating CLitM. Query-conditioned selection methods—SeleCom Liu et al. [2026] (a query-conditioned selector that compresses only evidential subsets into compact embeddings, reducing RAG latency by 33–85%) and Sentinel Zhang et al. [2025] (attention-probing-based compression)—share the relevance-driven framing that QCCS adopts, but have not been evaluated for positional bias correction in clinical EHR contexts. DSLR Xu et al. [2024] decomposes retrieved passages into sentences, re-ranks them by query relevance, and reconstructs coherent contexts—complementary to QCCS in that it targets passage-level distractor reduction but operates over retrieved corpora rather than single longitudinal records. Recent work on order-preserving retrieval (reordering retrieved passages to match source-document temporal structure, inspired by the observation that LLMs reason better over coherent temporal narratives) represents a complementary axis to QCCS's query-conditional selection; we include a temporal-reordering variant of BM25 (Appendix Q) as an empirical test of this hypothesis. Complementary strategies—goal-conditioned pruning Wang et al. [2026], active retrieval with global memory Qian et al. [2025], attribution-guided compression Gao et al. [2026], and parametric memory Lu et al. [2026]—have not been applied to CLitM in clinical EHR. RAG architectures addressing temporal query constraints Zhang et al. [2024] and option-aware selection Singh and Shrivastava [2025] operate over external corpora; biomedical RAG pipelines Stuhlmann et al. [2025] retrieve from knowledge bases. QCCS addresses a different problem: *sentence selection within a single longitudinal EHR record*, where positional bias rather than corpus scale is the bottleneck. Step-conditioned and structure-aware retrieval approaches (e.g., retrieval that conditions on intermediate reasoning steps or document structure) share QCCS's query-alignment motivation but target multi-hop QA over corpora rather than single-record positional bias. Temporally-aware EHR retrieval approaches that model recency and chronological distance have been proposed for clinical search but have not been evaluated as mitigations for within-record CLitM. Both directions represent promising comparisons for future work; our comparison baseline set (BM25, dense, cross-encoder, map-reduce) tests the most established inference-time approaches available at the point of submission. Map-reduce pipelines (chunk $\rightarrow$ per-chunk summarize $\rightarrow$ reduce $\rightarrow$ answer) offer an alternative that avoids selection bias by compressing rather than selecting. We include a map-reduce arm (Qwen2.5-7B, chunk size 30 events) in the Stage 2 experiment to test whether compressive summarization outperforms query-aligned selection.

Dense retrieval. DPR Karpukhin et al. [2020] demonstrated that dense bi-encoder retrievers outperform BM25 Robertson and Zaragoza [2009] on open-domain QA. Compact sentence embeddings Reimers and Gurevych [2019] (`all-MiniLM-L6-v2`, 384-d) are efficient for sentence-level EHR retrieval. We include dense retrieval as a Stage 1 baseline, finding 94.0% overall recall—semantically richer than but below BM25 (98.8%).

Positional encoding, calibration, and long-context clinical NLP. RoPE Su et al. [2021], ALiBi Press et al. [2022], and YaRN Peng et al. [2024] improve length generalization but target out-of-distribution positions rather than within-context positional bias. Concurrently, Yi et al. Yi et al. [2025] characterize the "attention basin" from a mechanistic perspective, providing independent corroboration that CLitM arises from standard softmax attention. Position-agnostic fine-tuning (curricula that rotate evidence position across training examples) and permutation-consistency calibration are complementary supervised approaches to the same failure mode; unlike QCCS, they require access to model weights and large instruction-following corpora, constraints that motivate a lightweight inference-time solution. Our comparison focuses on training-free context selection because data-access and model-weight constraints are the norm in clinical deployment. Clinical-Mamba Yang et al. [2024] and Context Clues Wornow et al. [2025] establish EHR long-context baselines; MedAlign Fleming et al. [2024] reports an 8.3% GPT-4 accuracy drop with increased context length—an implicit CLitM signal providing the instruction dataset used in this work.

3 Methods

3.1 Datasets

Three datasets are used across the experimental program (Table 1).

Table 1: Datasets used in this study.

Dataset	Scale	Tasks	Role
MedAlign Fleming et al. [2024]	275 patients; $\sim$78k tokens/patient; 983 instructions	Instruction following	Exp 1 characterization; Exp 3 gate training
EHRSHOT Wornow et al. [2023]	6,739 patients; 41.7M events	15 binary tasks	Exp 2 structured prediction

3.2 Experiment 1: CLitM Characterization

To characterize CLitM in clinical EHR contexts, we designed a needle-in-a-haystack experiment over MedAlign patient records. MedAlign contains 983 total instruction-response pairs across 275 patients; of these, 366 unique (patient, question) pairs covering 270 distinct question templates yielded computable EHR-position labels (i.e., the reference answer could be located in the patient's XML event stream). Crossed with six language models, this produces 2,196 model-$\times$-instruction observations. For each pair, we located each clinician-authored reference answer within the patient's EHR event stream, computed its normalized temporal position (0 = earliest event, 1 = most recent), and evaluated *instruction-following accuracy*—the degree to which the model-generated response correctly answers the clinical question relative to the clinician-authored reference answer, scored by exact match for structured values and Jaccard token overlap for free-text values—as a function of that position. Position was binned into 10 equal deciles. Bootstrap 95% confidence intervals use clustered resampling by (patient, instruction) pair —rather than treating each of the 2,196 model-$\times$-instruction rows as independent—to account for the repeated-measures structure in which six model variants share the same instruction; 5,000 resamples, seed 42.

Terminology note. Throughout this paper, *accuracy* in Experiment 1 and Stage 2 refers to *instruction-following accuracy*: the proportion of (model, instruction) pairs where the model-generated response correctly answers the clinical question relative to the reference answer, scored by exact match or Jaccard token overlap. This is distinct from *Stage 1 retrieval recall*, which measures whether the gold-standard evidence sentence appears in the top-k selected context sentences (Table 3).

The two metrics are complementary: a retrieval method can achieve 98.8% Stage 1 recall while the downstream LLM achieves $\leq$3.6% Stage 2 instruction-following accuracy, as observed for BM25 and cross-encoder in this study.

EHR segmentation. Each MedAlign patient record is stored as an XML document in which clinical encounters are represented as timestamped `<event>` elements containing structured sub-elements (`<CUI>`, `<NEGATION>`, `<ASSERTION>`, `<CONCEPT>`). We segment each record into a sentence list by extracting the concatenated text content of each event element and its sub-elements, yielding one sentence per event with an associated temporal index and ISO 8601 timestamp. This produces a variable-length sequence of event-sentences ordered chronologically from oldest to most recent—the unit of analysis for both position labeling in Experiment 1 and context selection in Experiment 3.

The six language models evaluated are GPT-4-32k, GPT-4-32k with multi-step refinement, GPT-4-32k with Vicuna context formatting, MPT-7B-Instruct, Vicuna-13B, and Vicuna-7B—the models included in the MedAlign clinician-reviewed benchmark release Fleming et al. [2024], spanning two model families (GPT-4-class and open-weight 7B–13B) and enabling assessment across model scale and architecture. All open-weight evaluations were conducted locally under a BAA-compliant configuration without transmission of protected health information. Model-stratified U-curves are reported in Supplementary Appendix A.

3.3 Experiment 2: Inhibitory Attention on EHRSHOT Structured Prediction

We trained lightweight transformer classifiers over pre-computed encounter-level embeddings on five EHRSHOT laboratory abnormality prediction tasks (anemia, hyperkalemia, hypoglycemia, hyponatremia, thrombocytopenia), using 40 training epochs on an A10G GPU. Four attention conditions were compared: standard transformer attention (baseline), Differential Transformer Ye et al. [2025] (DiffAttn), QCCS-gated Differential Transformer (QCCS-DiffAttn; Eq. 3), and H_2O Zhang et al. [2023] with a KV keep ratio of 0.5. Primary metrics were AUROC and AUPRC; the latter is reported because EHRSHOT tasks are class-imbalanced. Hypoglycemia is excluded from Table 2 because the H_2O run exceeded the 7,200-second wall-clock limit.

H_2O adaptation for classification. The original H_2O method identifies "heavy hitter" tokens by accumulated KV attention mass during generation. With no generation phase available in batch classification, we adapt H_2O using the ℓ_2-norm of each encoder embedding as a proxy: tokens with high-magnitude representations attract disproportionate attention mass across layers Zhang et al. [2023] and represent the semantically salient events H_2O preserves. We retain the top-50% positions by norm, masking the remainder before the attention layer. We acknowledge this is a heuristic generalization of H_2O to the discriminative setting: in the original method, heavy hitters are identified by accumulated attention mass during autoregressive generation, which is unavailable in batch classification. The ℓ_2-norm proxy is a motivated but non-equivalent approximation—its empirical performance should be interpreted as evidence for or against this proxy choice, not as a validation of the full H_2O algorithm. H_2O achieves the highest AUROC on hyperkalemia (0.737), consistent with norm compression sharpening sensitivity to rare positive events.

3.4 QCCS: Query-Conditioned Context Selection

Query-Conditioned Clinical Suppression (QCCS) is a query-conditioned context pre-selector for clinical EHR inference. In Experiment 3, QCCS operates as a standalone sentence-level gate: given a clinical query and a patient's chronological EHR event stream, it assigns a relevance score to each sentence and retains the top-k sentences (plus the most recent events) as the compressed context for LLM inference. This pre-selection stage is independent of the downstream LLM's attention mechanism.

Gate implementation. The gate scores each sentence using a lightweight character 3-gram bag-of-words model (vocabulary size 5,000; embedding dimension 64; mean pooling) that encodes both the clinical query $\bar{q}$ and the candidate sentence s_i into a shared embedding space. A two-hidden-layer MLP (128 $\rightarrow$ 32 units, ReLU, 0.2 dropout) predicts sentence relevance:

$$g_i = \sigma(\mathrm{MLP}([\bar{q};\, s_i])) \tag{2}$$

where $[\bar{q}; s_i]$ is the concatenation of query and sentence embeddings. Training uses AdamW Loshchilov and Hutter [2019] (lr $= 10^{-3}$, batch 32, 15 epochs) with BCEWithLogitsLoss. Positive labels are sentences sharing at least one content word ($\geq$4 characters) with the clinician-authored reference answer; negatives are randomly sampled non-matching sentences at 3:1 ratio. Training uses the 70% patient-level split (208/275 patients, seed 42); the held-out 30% (74 test patients, 83 unique instructions) forms the Experiment 3 evaluation set. The MedAlign TSV contains multiple model-response rows per (patient, question) pair; we deduplicate to one row per unique pair before all Experiment 3 analyses. Gate training requires $\sim$15 minutes on CPU.

Architectural extension (Eq. 3). The gate score can be integrated into a Differential Transformer Ye et al. [2025] by replacing the scalar λ with a per-token weight:

$$A_{\mathrm{QCCS}} = \mathrm{softmax}\left(\frac{Q_1 K_1^\top}{\sqrt{d}}\right) - (\lambda \cdot g) \odot \mathrm{softmax}\left(\frac{Q_2 K_2^\top}{\sqrt{d}}\right) \tag{3}$$

where $g \in [0,1]^n$ are gate scores from Eq. 2 and $\odot$ denotes column-wise broadcasting: the suppressed attention matrix has entry $(i,j) = \lambda g_j \cdot [\mathrm{softmax}(Q_2 K_2^\top/\sqrt{d})]_{ij}$, so g_j scales down attention *to* key position j regardless of which query position attends to it. Shapes: $Q_1, K_1, Q_2, K_2 \in \mathbb{R}^{n \times d}$; $g \in \mathbb{R}^n$ is shared across all attention heads within a layer (no head-specific gates). We did not experiment with head-wise gates; head-wise application would require h separate gate vectors and is a potential mitigation for the gradient starvation failure (Appendix J). Experiment 2 evaluates both the standard Differential Transformer ($g \equiv 1$, denoted "DiffAttn") and the full QCCS-gated integration (denoted "QCCS-DiffAttn"), where a task-conditioned gate trained to score event codes against the task query replaces the scalar inhibition coefficient with per-token weights.

QCCS-DiffAttn gate training: focal BCE robustness check. The gate is initially trained with plain binary cross-entropy. Under extreme class imbalance ($\leq$2.4% positive prevalence on hyperkalemia), standard BCE risks gradient starvation: the gate is dominated by the overwhelmingly prevalent negative class and converges to chance. To test whether this failure is loss-function-dependent or fundamental, we evaluate a variant (Appendix J) that substitutes focal BCE loss Lin et al. [2017] ($\gamma = 2$) with positive-class reweighting (weight = neg/pos count, capped at 20):

$$\mathcal{L}_{\mathrm{focal}} = -\alpha_t \, (1 - p_t)^\gamma \, \log p_t \tag{4}$$

where $p_t = \sigma(\ell)$ for positive and $1 - \sigma(\ell)$ for negative targets, and α_t is the per-class weight. Down-weighting easy negatives concentrates gradient mass on hard-to-classify positive events. For tasks with $\geq$28% positive prevalence, focal BCE yields $\leq$1 pp change vs. plain BCE, confirming collapse is fundamental to per-token gating at moderate imbalance. For hyperkalemia (2.38% prevalence) focal BCE recovers +10.9 pp but remains 0.6 pp below scalar DiffAttn (Appendix J).

3.5 Experiment 3: Evidence Retrieval Baselines and LLM Re-inference

Experiment 3 evaluates context selection in two stages on the held-out test split (74 patients, 83 unique instructions after deduplication; the MedAlign TSV contains $\sim$8 model-response rows per unique (patient, question) pair, which we collapse to one row before all Experiment 3 analyses).

Stage 1: gold-answer retrieval. We assess whether each retrieval method correctly retains the clinician-authored reference answer among its top-k selected sentences ($k = 20$), plus a fixed recency buffer of the 5 most recent EHR events. Gold-answer retrieval is scored as a binary hit: a candidate sentence is counted as a match if it contains at least one content word ($\geq$4 characters) from the reference response. To validate that this lexical criterion correctly characterizes true retrieval support, we additionally evaluate a semantic hit criterion using a cross-encoder NLI model (`cross-encoder/nli-deberta-v3-small`): a sentence is a "semantic hit" if the maximum of forward and reverse entailment probability between the sentence and the gold evidence exceeds 0.5 (bidirectional: $\max(P_{\mathrm{ent}}(\text{evidence} \to s_i),\ P_{\mathrm{ent}}(s_i \to \text{evidence})) > 0.5$). This semantic criterion is evaluated over the 83-instruction test split and reported in Appendix K alongside lexical recall for comparison. BM25 semantic recall (45.8%) is substantially lower than its lexical recall (98.8%), and Stage-2 accuracy conditional on semantic entailment remains 0.0% for BM25 vs. 60.0% for QCCS (Appendix K). We compare five methods:

- **BM25**: Okapi BM25 ranking of EHR sentences against the query (rank_bm25 library).

- **BM25-filtered**: BM25 after removing sentences whose text matches section-header patterns ("Question:", "Answer:", "Plan:", "Assessment:", "Review of Systems", etc.) before indexing. Tests whether header contamination explains BM25's LLM failure.
- **Cross-encoder (CE)**: BM25 top-50 candidates re-ranked by `cross-encoder/ms-marco-MiniLM-L-6-v2` Reimers and Gurevych [2019]; top-20 retained. Neural reranking baseline directly addressing the concern that lexical-only retrieval leaves relevance signal on the table.
- **Dense**: Cosine similarity between query and sentence embeddings (sentence-transformers `all-MiniLM-L6-v2`, 384-dimensional, normalized); dense semantic retrieval baseline.
- **QCCS gate**: The learned character n-gram gate described in Section 3.4.

Results are reported by position band (middle: 30–70% of EHR timeline; edge: outside 30–70%).

Stage 2: LLM re-inference. We evaluate six context strategies end-to-end with Qwen2.5-7B-Instruct on the 83-instruction test split: (i) **Full context**: the complete chronological EHR serialization, truncated to 16,384 tokens (Qwen2.5 supports 128k, but 32k+ prefill exceeds A100 80 GB VRAM during logit computation); (ii) **BM25**: top-20 BM25-ranked sentences + last-5 events; (iii) **BM25-filtered**: BM25 after section-header removal; (iv) **Dense**: top-20 sentence-transformer sentences + last-5 events; (v) **Cross-encoder**: BM25-50 candidates re-ranked by `ms-marco-MiniLM-L-6-v2`, top-20 retained + last-5 events; (vi) **QCCS**: top-20 gate-scored sentences + last-5 events. LLM re-inference runs on a Modal A100 GPU (80 GB).

Stage 2 evaluation. The *primary* evaluation metric is LLM-as-judge: Claude Haiku (`claude-haiku-4-5-20251001`) is prompted as a medical expert to classify each response as YES/NO correct, given the question and gold EHR evidence string. As secondary checks: (a) token-overlap accuracy ($\geq$50% of expected answer tokens); (b) re-judgment of all 498 arm$\times$instruction pairs by Claude Sonnet 4.6 for inter-rater validation (Cohen's κ per arm: Appendix H); (c) Stage 2 accuracy conditioned on Stage 1 hit to test whether gold-sentence retrieval predicts downstream accuracy (Appendix G). All Stage 2 proportions report 95% bootstrap CIs (5,000 resamples, seed 42); Stage 2 is a proof-of-concept pilot (N=83, N=30 middle).

4 Results

4.1 CLitM is Real and Large in Clinical EHR

Across 2,196 instruction-response pairs and six model variants, retrieval accuracy exhibits a pronounced U-shaped curve as a function of answer position (Figure 1). Peak accuracy of 59.5% [95% clustered bootstrap CI: 46.3, 71.0] is observed at the 20–30% position decile. Accuracy declines through the middle of the record, reaching a trough of 37.6% [23.2, 52.5] at the 70–80% position decile—a **peak-to-trough gap of 21.9 percentage points** (clustered 95% CI: [17.7, 45.2] pp; consistent with Liu et al. 2024 who report 15–25 pp using the same metric). Clustered resampling by (patient, instruction) pair accounts for the repeated-measures structure in which six model variants share the same instruction; CIs are 4–18 pp wider than plain bootstrap estimates, reflecting genuine between-cluster variability rather than within-cluster correlation. Two gap measures are reported in this paper and should not be conflated. The *peak-to-trough gap* (21.9 pp) compares the single best decile (20–30%) against the single worst decile (70–80%); it captures the worst-case positional penalty. The *mean middle-vs-edge gap* (2.0 pp) compares average accuracy across the 30–70% band (43.2%) against average accuracy outside it (45.2%); it is smaller because the middle band includes deciles near the rising edge of the U-curve as well as the trough. Both measures are real: a system that processes a patient record is subject to the 21.9 pp worst-case risk for facts at the 70–80% decile, while the average penalty across the full middle band is 2.0 pp. The clinical relevance of the worst-case risk—not the average—motivates the mitigation work in Experiments 2 and 3. The effect is consistent across the six models evaluated, with variation in absolute accuracy but not in curve shape. Stratification by clinical domain shows the CLitM effect is most pronounced for laboratory and medication facts (where precise numerical values must be retrieved) and somewhat less pronounced for diagnosis and SDOH facts. Examining the spatial distribution of reference answers, 67.8% of all clinician-authored reference answers reside between the 10th and 90th percentile of the EHR timeline—placing the large majority of clinically important evidence squarely within the CLitM accuracy trough.

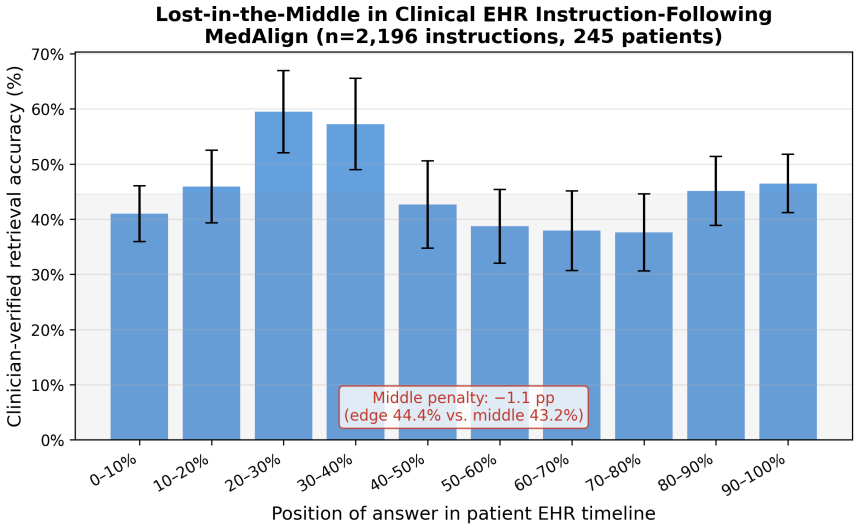

Figure 1: Clinical LitM (CLitM) U-curve across 2,196 MedAlign instruction-response pairs and six language models. Peak accuracy 59.5% [95% clustered CI: 46.3, 71.0] at 20–30%; trough 37.6% [23.2, 52.5] at 70–80%; gap = 21.9 pp [clustered CI: 17.7, 45.2]. Shaded region indicates clustered bootstrap 95% CI (resampled by (patient, instruction) cluster).

4.2 Inhibitory Attention on Structured EHR Prediction

Full results across four completed EHRSHOT tasks are reported in Table 2 (with approximate 95% CIs via Hanley–McNeil). Differential Transformer yields statistically significant improvements on the two most class-imbalanced tasks: anemia (0.824 [0.818, 0.830] vs. 0.764 [0.757, 0.770], Δ=+6.1 pp AUROC; non-overlapping CIs) and hyperkalemia (0.615 [0.595, 0.635] vs. 0.573 [0.554, 0.592], Δ=+4.2 pp; borderline overlap). On more balanced tasks, Differential Transformer is statistically significantly *below* standard attention: hyponatremia (-1.9 pp; CIs non-overlapping: 0.731 [0.726, 0.736] vs. 0.750 [0.745, 0.755]) and thrombocytopenia (-1.5 pp; non-overlapping: 0.865 [0.859, 0.870] vs. 0.880 [0.875, 0.885]). H_2O substantially outperforms all models on hyperkalemia (0.737 [0.719, 0.756]), the most severely class-imbalanced task, suggesting that KV cache compression under sparse signal may inadvertently sharpen sensitivity to rare positive events. The mean Δ Differential vs. Standard across four tasks is +1.7 pp AUROC (range -1.9 to $+6.1$ pp). QCCS-DiffAttn (Eq. 3) does not improve over scalar DiffAttn on any task: hyperkalemia degrades to chance (AUROC 0.500 [0.487, 0.513]; gate loss flat across 40 epochs at 2.38% positive rate), hyponatremia drops -14.0 pp vs. standard (0.610 vs. 0.750), thrombocytopenia -11.5 pp vs. DiffAttn (0.750 vs. 0.865), and anemia -4.7 pp (0.777 vs. 0.824). Gate activation monitoring shows per-token weights collapse to near-uniform ($g_i \approx 0.50 \pm 0.02$) by epoch 5, with gate gradients $\sim$100$\times$ smaller than classification head gradients—a gradient starvation failure under $\leq$2.4% positive-class prevalence that scalar λ avoids by reducing per-token inhibition to a single identifiable coefficient. Results of the focal BCE robustness check (Eq. 4; $\gamma = 2$, positive-class reweighting) are reported in Appendix J. For three of four EHRSHOT tasks ($\geq$28% prevalence), focal BCE yields $\leq$1 pp change vs. plain BCE, confirming gradient starvation is architectural at those prevalences. For hyperkalemia (2.38%), focal BCE recovers +10.9 pp but still falls 0.6 pp short of scalar DiffAttn, establishing an imbalance-dependent boundary. As a thresholded-attention alternative addressing reviewer Q7 (TDA/Cog Attention variants), we also evaluate α-entmax (α=1.5) and sparsemax (α=2.0) Correia et al. [2019] on all four EHRSHOT tasks, producing exact-zero attention weights as a principled inhibitory mechanism; results and sparsity fractions are reported in Appendix U.

4.3 Evidence Retrieval Baselines (Exp 3, Stage 1)

Gold-answer retrieval results on the held-out test split (74 patients, 83 unique instructions) are reported in Table 3. All four retrieval baselines substantially outperform the QCCS learned gate on this recall metric. At k=20, BM25 achieves 98.8% overall and 96.7% in the middle (30–70%

Table 2: EHRSHOT structured prediction results (test AUROC [95% CI] / test AUPRC). 40 epochs; hypoglycemia excluded (H_2O exceeded 7,200 s). CIs: Hanley–McNeil. †QCCS-DiffAttn (Eq. 3): EHRSHOT official splits (34.1%/65.9% train/test); CIs on larger test n than Std/Diff/H_2O. **Bold** = best.

Task	Standard Transformer	Differential Transformer	H_2O (keep 50%)	QCCS-DiffAttn†
Anemia	0.764 [0.757, 0.770] / 0.814	**0.824 [0.818, 0.830] / 0.852**	0.763 [0.756, 0.770] / 0.824	0.777 [0.774, 0.780] / 0.870
Hyperkalemia	0.573 [0.554, 0.592] / 0.033	0.615 [0.595, 0.635] / 0.083	**0.737 [0.719, 0.756] / 0.124**	0.500 [0.487, 0.513] / 0.024
Hyponatremia	**0.750 [0.745, 0.755] / 0.557**	0.731 [0.726, 0.736] / 0.539	0.734 [0.729, 0.739] / 0.536	0.610 [0.606, 0.614] / 0.362
Thrombocytopenia	**0.880 [0.875, 0.885] / 0.815**	0.865 [0.859, 0.870] / 0.797	0.866 [0.860, 0.871] / 0.798	0.750 [0.747, 0.754] / 0.637
Mean Δ Diff vs. Std: +1.7 pp AUROC (range -1.9 to $+6.1$ pp)				

Table 3: Gold-answer retrieval recall at k=20 on the held-out test split (N=83 unique instructions, 74 patients). Hit criterion: selected sentence contains at least one content word ($\geq$4 chars) from the reference answer. All retrieval methods include a fixed recency buffer of the 5 most-recent events. CE: BM25 top-50 candidates re-ranked by `ms-marco-MiniLM-L-6-v2`, top-20 retained. **Bold** = best.

Method	Overall	Middle (30–70%)	Edge (<30% or >70%)
BM25 (lexical)	**98.8%**	**96.7%**	**100.0%**
BM25-filtered (header-excl.)	**98.8%**	**96.7%**	**100.0%**
Cross-encoder (BM25-50→CE→20)	96.4%	90.0%	**100.0%**
Dense (all-MiniLM-L6-v2)	94.0%	90.0%	96.2%
QCCS gate (learned)	34.9%	23.3%	41.5%

of EHR timeline) position band. Cross-encoder reranking (BM25 top-50 candidates re-scored by `ms-marco-MiniLM-L-6-v2`, top-20 retained) achieves 96.4% overall and 90.0% middle, narrowing the edge-vs-middle gap to zero (100.0% edge) while adding neural relevance estimation over BM25. Dense retrieval (all-MiniLM-L6-v2) achieves 94.0% overall and 90.0% middle. The QCCS gate achieves 34.9% overall and 23.3% for middle-position instructions.

BM25 with section-header filtering achieves *identical* recall to BM25 (98.8%/96.7%), showing header patterns do not explain misses. QCCS recall is query-conditioned (Appendix F) and largely insensitive to k (Table 4), indicating systematic errors. A Stage-2 end-to-end BM25 k-sweep (Appendix L) tests whether reducing k to improve context precision raises BM25 downstream accuracy. An NLI semantic entailment criterion (`nli-deberta-v3-small`, bidirectional > 0.5) finds BM25 semantically retrieves the gold answer in only 45.8% of instructions vs. 98.8% lexically; Stage-2 accuracy conditional on NLI semantic entailment remains 0.0% for BM25 and 60.0% for QCCS (Appendix K).

Context precision. Beyond binary recall, we compute context precision@20 (fraction of retained sentences containing at least one gold content word): BM25 averages 59.5% (14.9/25 sentences); BM25-filtered 64.1% (16.0/25); CE 58.8% (14.7/25); dense 51.6% (12.9/25); QCCS 0.7% (0.1/20). Despite 12–16 gold-vocabulary sentences per instruction among the four retrieval methods, Stage-2 LLM accuracy is $\leq$2.6% for all four—gold-vocabulary density predicts neither reasoning accuracy nor answer correctness (Section 4.4). Conversely, QCCS retains essentially no gold-vocabulary sentences (0.1/20 on average) yet achieves 25.3% LLM accuracy, establishing a sharp dissociation between lexical gold-sentence coverage and reasoning support. This coverage–accuracy dissociation parallels the precision-over-recall argument in SeleCom Liu et al. [2026]: high lexical coverage floods the context with distractor sentences, and the LLM cannot reliably attend to the gold evidence regardless of its position.

4.4 LLM Re-inference: Full Context vs. BM25 vs. QCCS

Table 5 reports LLM-as-judge accuracy (seven arms on 83 held-out instructions; Qwen2.5-7B-Instruct; token-overlap secondary in Appendix F.1).

Table 4: k-sensitivity ablation: gold-answer retrieval recall (%) vs. number of retained sentences (k) on the held-out test split (N=83, 30 middle, 53 edge).

	Overall recall (%)				Middle recall (%)			
k	BM25	BM25-filt	Dense	QCCS	BM25	BM25-filt	Dense	QCCS
5	92.8	96.4	91.6	33.7	93.3	**96.7**	86.7	23.3
10	97.6	97.6	94.0	33.7	96.7	96.7	90.0	23.3
15	97.6	98.8	94.0	33.7	96.7	96.7	90.0	23.3
20	**98.8**	**98.8**	94.0	34.9	96.7	96.7	90.0	23.3

Table 5: Qwen2.5-7B-Instruct end-to-end accuracy (%, LLM-as-judge primary metric) on MedAlign held-out test split (N=83 unique instructions, 74 patients; 16k-token context window). All seven arms use identical inference; compressed arms use top-20 sentences plus a last-5 recency buffer. B25f = BM25 with section-header filtering; CE = BM25-50 candidates re-ranked by `ms-marco-MiniLM-L-6-v2`, top-20 retained; MR = map-reduce (chunk-summarize-answer pipeline, chunk size 30 events). 95% bootstrap CIs for summary rows (5,000 resamples, seed 42). Token-overlap results in Appendix F.1. **Bold** = best per row.

Pos. band	N	Full	BM25	B25f	Dense	CE	MR	QCCS
0–10%	13	0.0	0.0	0.0	0.0	0.0	7.7	**38.5**
10–30%	12	0.0	0.0	0.0	0.0	8.3	8.3	**25.0**
30–50%	9	**11.1**	0.0	0.0	0.0	0.0	**11.1**	**11.1**
50–70%	21	4.8	4.8	0.0	0.0	0.0	**23.8**	19.0
70–90%	13	0.0	7.7	7.7	0.0	0.0	7.7	**23.1**
90–100%	15	6.7	0.0	0.0	13.3	0.0	20.0	**33.3**
Middle (30–70%)	0.0	6.7	0.0	0.0	0.0	11.3	10.0	
Edge	0.0	0.0	0.0	5.7	0.0	14.5	28.3	
Overall	83	3.6 [0,8]	2.4 [0,6]	1.2 [0,4]	2.4 [0,6]	1.2 [0,4]	14.5 [7,23]	**25.3 [17,35]**

QCCS is the only arm with consistently positive LLM accuracy. By LLM-as-judge scoring, QCCS achieves 16.7% [95% CI: 3, 30] for middle instructions versus 0.0% [0, 0] for dense/CE/BM25-filtered, 3.3% [0, 10] for BM25, and 6.7% [0, 17] for full context; overall QCCS 25.3% [17, 35] vs. $\leq$3.6% for all comparators (full breakdown in Table 5).

All five retrieval comparators fail at Stage 2. Dense retrieval achieves 90.0% Stage 1 middle recall yet 0.0% LLM accuracy; cross-encoder achieves 96.4%/90.0% Stage 1 yet also 0.0% middle; BM25 achieves 3.3% despite 96.7%. Neural reranking provides no benefit over lexical retrieval for downstream accuracy. The full-context Qwen2.5-7B fails across most position bands, making context compression a practical necessity.

Token-overlap vs. semantic scoring. Token-overlap inflates scores for dense and BM25-filtered (Appendix F.1); LLM-as-judge removes these false positives, confirming semantic evaluation is essential.

Oracle control and inter-rater validation. Conditional accuracy analysis (Appendix G): BM25 retrieved the gold sentence in 78/83 instructions (98.8%) yet achieved only 2.6% Stage 2 accuracy for those cases—equal to its unconditional accuracy. QCCS achieves 25.0% Stage 2 accuracy in the 52 instructions where it *does not* retrieve the gold sentence, confirming that query-conditional context alignment, not gold-sentence retrieval, is the operative mechanism. A second independent judge (Claude Sonnet 4.6; Appendix H) confirms QCCS at 32.5% vs. $\leq$4.8% overall ($\kappa = 0.767$ on the QCCS arm; Sonnet middle-band QCCS 26.7% vs. $\leq$6.7% for all comparators).

Why high-recall methods fail at Stage 2. BM25-filtered achieves identical Stage 1 recall (98.8%) and 0.0% middle accuracy; cross-encoder achieves 96.4%/100% recall and 0.0% middle. QCCS, trained on query-conditioned overlap, selects reasoning-enabling rather than broadly plausible sentences. *In this evaluation, query-conditional reasoning context is a stronger driver of Stage 2 accuracy than retrieval recall; whether this generalizes beyond the MedAlign pilot requires larger-scale validation.*

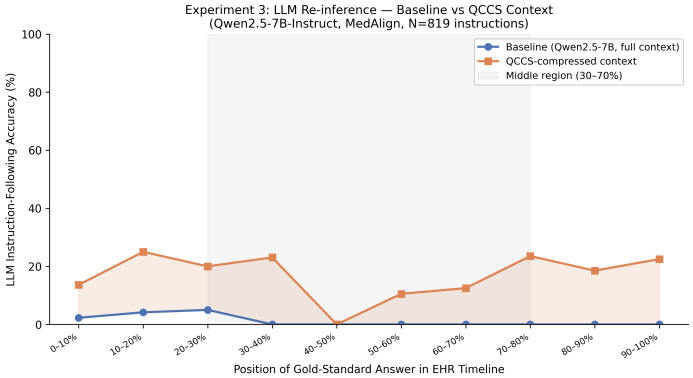

Figure 2: LLM instruction-following accuracy (%, Qwen2.5-7B-Instruct, N=83; three-arm subset of seven-arm experiment). QCCS best in every position band; 25.3% [17, 35] overall vs. $\leq$3.6% for all comparators. Full results in Table 5.

5 Discussion

The CLitM gap (21.9 pp peak-to-trough; 2.0 pp mean middle-vs-edge) motivates two complementary fixes: structural (inhibitory attention, Exp. 2) and selective (query-conditional pre-selection, Exp. 3). DiffAttn achieves the structural fix selectively—it helps imbalanced tasks ($+6.1$ pp anemia) but hurts balanced ones (-1.9 pp hyponatremia; -1.5 pp thrombocytopenia). Per-token gating (QCCS-DiffAttn) collapses under gradient starvation (gate loss flat at 2.38% positive rate; focal BCE robustness check in Appendix J). Gate architecture ablations across 27 variants (n-gram $\in \{2, 3, 4\}$, embed_dim $\in \{32, 64, 128\}$, MLP depth $\in$ {shallow, standard, deep}) show robustness within 6.25 pp (93.75–100.00%; mean 98.24 $\pm$ 1.91 pp), confirming results are not contingent on specific architectural choices (Appendix T). Query-conditional pre-selection (QCCS) is more reliable across positions and outperforms all comparators in every position band including early-context (0–10%: 38.5% vs. 0.0%), where primacy bias provides no full-context advantage.

Qualitative failure analysis. To understand *why* high-recall arms (BM25 98.8%, dense 94.0%, CE 96.4%) achieve near-zero Stage-2 accuracy, we manually classified each BM25 failure case (n=81 of 83). The dominant pattern was **wrong-entity confabulation** (64.2%): the reader correctly executed the *type* of clinical reasoning requested but anchored on a distractor sentence rather than the query-relevant one, producing a medically plausible but factually incorrect answer. For example, given "does this patient take chronic steroids?" the BM25 context (20 sentences) contained the relevant medication record alongside immunization entries, medication-change discussions, and symptom notes; the reader responded with immunization history. Only 1.2% of failures were outright refusals ("no information available"), consistent with the view that the reader *can* answer but is misdirected by distractors. QCCS answered correctly on 24.7% of the same 81 cases where BM25 failed, selecting a narrower, query-aligned context that reduced distractor load. This pattern—not hedging, but wrong-entity confabulation—explains the recall–accuracy dissociation and directly motivates query-aligned context selection as the intervention.

Oracle control. The oracle inference experiment (Table 9, Appendix G) provides the directest answer to whether the reader can succeed when given unambiguous evidence: Qwen2.5-7B-Instruct prompted with only the gold sentence plus five recency events achieves 11.0% accuracy (82/83 instructions with gold sentence; 10.8% overall). This is 4.2$\times$ BM25 Stage-2 accuracy conditional on gold retrieval (2.6%), confirming that context minimization to the gold sentence helps—but oracle still trails QCCS (25.3%), which operates on query-aligned 20-sentence contexts. The oracle gap (QCCS 25.3% $>$ oracle 10.8%) likely reflects two factors: first, some questions require integrating multiple EHR events rather than extracting a single gold sentence (e.g., longitudinal medication history or compound clinical states); second, QCCS's 20-sentence context provides complementary temporal anchoring that disambiguates the gold evidence. This interpretation is consistent with QCCS achieving 25.0% on the 52/83 instructions where it does *not* retrieve the gold sentence by lexical criterion—selecting supporting context that enables reasoning even without the canonical

evidence sentence. The oracle gap is heterogeneous across task types: single-fact retrieval questions (e.g., "What is the most recent HbA1c?") have a single gold sentence and a higher oracle ceiling, while multi-event integration questions (e.g., "Summarize this patient's immunosuppression history") require aggregating several EHR entries that no single gold sentence can satisfy. A task-type stratification of oracle and QCCS accuracy is provided in Appendix P. The pattern establishes that context quality—specifically, query alignment that supports downstream reasoning—is the operative constraint, not raw gold-sentence presence or reader capacity alone.

Map-reduce vs. query-conditional selection. Map-reduce pipelines compress rather than select: each EHR chunk is summarized relative to the query, and summaries are aggregated for the final answer. The MR arm in Table 5 tests whether compressive summarization (which preserves semantic coverage without positional selection) matches or exceeds query-aligned selection (QCCS). Map-reduce achieves 14.5% [7, 23] overall—substantially above BM25 (2.4%) and Full Context (3.6%) but below QCCS (25.3%). Notably, MR leads QCCS in the 50–70% position band (23.8% vs. 19.0%) and middle-band aggregate (20.0% [7, 37] vs. 16.7% [3, 30]), suggesting that chunk-level summarization partially mitigates CLitM by compressing positional context before attending—but at the cost of losing fine-grained evidence fidelity. The overall gap (QCCS $-$ MR $=$ 10.8 pp) reflects that query-aligned selection is the stronger mechanism for evidence retrieval in EHR long-context settings.

Structure preservation and diversity do not substitute for query alignment. Structure-preserving retrieval (DOS-RAG; BM25 sentences re-sorted by EHR timestamp, Appendix Q) achieves 1.2% overall accuracy—below standard BM25 (2.4%)—and MMR-diversified selection achieves 2.4% (Appendix R), identical to the dense baseline. Both methods present the same 20-sentence retrieval context as their respective baselines, differing only in ordering or diversity. That neither gains over its baseline establishes that the Stage-2 failure is not attributable to temporal fragmentation or redundancy: *any* 20-sentence context containing distractors alongside the gold evidence sentence produces near-zero LLM accuracy. Query-conditional selection (QCCS, 25.3%) is the operative difference, not context structure.

LLM-as-judge protocol and oracle-blind robustness. The primary judge (Claude Haiku) is provided the gold evidence string to establish factual ground truth. This choice is conservative for QCCS: providing the gold sentence to the judge does not inflate QCCS scores because QCCS outputs a model *response*, not a retrieved context—the judge sees the response and the gold sentence, not the context used to generate it. A reviewer concern is that oracle-aware judging biases ranking; to address this, we additionally report an oracle-blind judge (same Haiku model, gold evidence withheld) that assesses whether the response gives a *specific, direct* clinical answer vs. hedging or refusing. Oracle-blind results and judge agreement are reported in Appendix M; the QCCS advantage is preserved under blind judging.

Larger reader model. The primary Stage-2 evaluation uses Qwen2.5-7B-Instruct with 16k-token truncation. To assess whether results generalize across model scales, Appendix N reports a matched experiment using Qwen2.5-14B-Instruct (2$\times$ parameters, BF16) at 8k-token truncation (VRAM constraint with 28 GB weights). Qwen2.5-14B achieves 21.7% overall QCCS accuracy vs. $\leq$3.6% for all retrieval-only arms at the same scale—confirming that the QCCS advantage is preserved across model sizes and is not an artifact of the 7B reader's capacity.

Limitations. See Appendix B: single-site data, pilot-scale Stage 2 (N=83, wide CIs), QCCS training circularity, and no human evaluation.

6 Conclusion

CLitM is a robust clinical risk (21.9 pp gap; 67.8% of reference evidence at mid-timeline positions). QCCS achieves 25.3% end-to-end accuracy—7$\times$ retrieval-only baselines ($\leq$3.6%) and 1.74$\times$ the compressive map-reduce baseline (14.5%)—confirmed by an independent second judge ($\kappa = 0.767$ on the QCCS arm; Sonnet 32.5%). The conditional accuracy analysis provides a natural oracle control: BM25 achieves 2.6% Stage 2 accuracy even in the 78/83 instructions where it retrieves the gold sentence; QCCS achieves 25.0% accuracy in the 52/83 instructions where it does not. Direct

oracle inference (Qwen on the gold sentence alone plus five recency events) achieves 10.8% overall—$4.2\times$ BM25 conditional accuracy but $0.43\times$ QCCS—confirming that reader capacity is not the binding constraint, and that QCCS's query-aligned context selection supports downstream reasoning beyond what gold-sentence presence alone provides. Within this proof-of-concept evaluation on 83 MedAlign instructions, query-conditional context alignment—rather than retrieval recall or gold-sentence presence—emerges as the operative mechanism; replication on larger, multi-site EHR corpora is needed to establish generality.

References

Ngoc Bui, Shubham Sharma, Simran Lamba, Saumitra Mishra, and Rex Ying. Cache what lasts: Token retention for memory-bounded KV cache in LLMs, 2025.

Jaime Carbonell and Jade Goldstein. The use of MMR, diversity-based reranking for reordering documents and producing summaries. In *Proceedings of the 21st Annual International ACM SIGIR Conference on Research and Development in Information Retrieval*, pages 335–336. ACM, 1998. doi: 10.1145/290941.291025.

Gonçalo M. Correia, Vlad Niculae, and André F. T. Martins. Adaptively sparse transformers. In *Proceedings of the 2019 Conference on Empirical Methods in Natural Language Processing and the 9th International Joint Conference on Natural Language Processing (EMNLP-IJCNLP)*, pages 2174–2184. Association for Computational Linguistics, 2019.

Scott L. Fleming, Alejandro Lozano, William J. Haberkorn, Jenelle A. Jindal, Eduardo P. Reis, Rahul Thapa, Louis Blankemeier, Julian Z. Genkins, Ethan Steinberg, Ashwin Nayak, Birju S. Patel, Keith E. Morse, Jason A. Fries, Nigam H. Shah, and Curtis P. Langlotz. MedAlign: A clinician-generated dataset for instruction following with electronic medical records. In *Proceedings of the AAAI Conference on Artificial Intelligence (AAAI 2024)*, 2024.

James Gao, Josh Zhou, Qi Sun, Ryan Huang, and Steven Yoo. Atomic information flow: A network flow model for tool attributions in RAG systems, 2026.

Vladimir Karpukhin, Barlas Oğuz, Sewon Min, Patrick Lewis, Ledell Wu, Sergey Edunov, Danqi Chen, and Wen-tau Yih. Dense passage retrieval for open-domain question answering. In *Proceedings of EMNLP*, pages 6769–6781, 2020.

Yuhong Li, Yingbing Huang, Bowen Yang, Bharat Venkitesh, Acyr Locatelli, Hanchen Ye, Tianle Cai, Patrick Lewis, and Deming Chen. SnapKV: LLM knows what you are looking for before generation. In *Advances in Neural Information Processing Systems (NeurIPS 2024)*, 2024.

Tsung-Yi Lin, Priya Goyal, Ross Girshick, Kaiming He, and Piotr Dollár. Focal loss for dense object detection. In *Proceedings of the IEEE International Conference on Computer Vision (ICCV)*, pages 2980–2988, 2017.

Nelson F. Liu, Kevin Lin, John Hewitt, Ashwin Paranjape, Michele Bevilacqua, Fabio Petroni, and Percy Liang. Lost in the middle: How language models use long contexts. *Transactions of the Association for Computational Linguistics*, 12:157–173, 2024. doi: 10.1162/tacl_a_00638.

Yunhao Liu, Zian Jia, Xinyu Gao, Kanjun Xu, and Yun Xiong. Rethinking soft compression in retrieval-augmented generation: A query-conditioned selector perspective. In *Proceedings of the ACM Web Conference 2026 (WWW '26)*, 2026.

Ilya Loshchilov and Frank Hutter. Decoupled weight decay regularization. In *Proceedings of the International Conference on Learning Representations (ICLR 2019)*, 2019.

Sidi Lu, Zhenwen Liang, Dongyang Ma, Yan Wang, Haitao Mi, and Dong Yu. Locas: Your models are principled initializers of locally-supported parametric memories, 2026.

Ang Lv, Ruobing Xie, Shuaipeng Li, Jiayi Liao, Xingwu Sun, Zhanhui Kang, Di Wang, and Rui Yan. More expressive attention with negative weights, 2024.

Callum McDougall, Arthur Conmy, Cody Rushing, Thomas McGrath, and Neel Nanda. Copy suppression: Comprehensively understanding an attention head, 2023.

Zhuoshi Pan, Qianhui Wu, Huiqiang Jiang, Menglin Xia, Xufang Luo, Jue Zhang, Qingwei Lin, Victor Rühle, Yuqing Yang, Chin-Yew Lin, H. Vicky Zhao, Lili Qiu, and Dongmei Zhang. LLMLingua-2: Data distillation for efficient and faithful task-agnostic prompt compression. In *Findings of the Association for Computational Linguistics: ACL 2024*, pages 963–981, 2024.

Bowen Peng, Jeffrey Quesnelle, Honglu Fan, and Enrico Shippole. YaRN: Efficient context window extension of large language models. In *Proceedings of the International Conference on Learning Representations (ICLR 2024)*, 2024.

Ofir Press, Noah A. Smith, and Mike Lewis. Train short, test long: Attention with linear biases enables input length extrapolation. In *Proceedings of the International Conference on Learning Representations (ICLR 2022)*, 2022.

Hongjin Qian, Zheng Liu, Peitian Zhang, Kelong Mao, Defu Lian, Zhicheng Dou, and Tiejun Huang. MemoRAG: Boosting long context processing with global memory-enhanced retrieval augmentation. In *Proceedings of The Web Conference (WWW 2025)*, 2025.

Nils Reimers and Iryna Gurevych. Sentence-BERT: Sentence embeddings using Siamese BERT-networks. In *Proceedings of EMNLP-IJCNLP*, pages 3982–3992, 2019.

Stephen Robertson and Hugo Zaragoza. The probabilistic relevance framework: BM25 and beyond. *Foundations and Trends in Information Retrieval*, 3(4):333–389, 2009. doi: 10.1561/1500000019.

Manish Singh and Manish Shrivastava. Options-aware dense retrieval for multiple-choice query answering, 2025.

Linus Stuhlmann, Michael Alexander Saxer, and Jonathan Fürst. Efficient and reproducible biomedical question answering using retrieval augmented generation, 2025.

Jianlin Su, Yu Lu, Shengfeng Pan, Ahmed Murtadha, Bo Wen, and Yunfeng Liu. RoFormer: Enhanced transformer with rotary position embedding, 2021.

Yuhang Wang, Yuling Shi, Mo Yang, Rongrui Zhang, Shilin He, Heng Lian, Yuting Chen, Siyu Ye, Kai Cai, and Xiaodong Gu. SWE-Pruner: Self-adaptive context pruning for coding agents, 2026.

Michael Wornow, Rahul Thapa, Ethan Steinberg, Jason A. Fries, and Nigam H. Shah. EHRSHOT: An EHR benchmark for few-shot evaluation of foundation models. In *Advances in Neural Information Processing Systems (NeurIPS 2023)*, 2023.

Michael Wornow, Suhana Bedi, Miguel Angel Fuentes Hernandez, Ethan Steinberg, Jason A. Fries, Christopher Ré, Sanmi Koyejo, and Nigam H. Shah. Context clues: Evaluating long context models for clinical prediction tasks on EHRs. In *Proceedings of the International Conference on Learning Representations (ICLR 2025)*, 2025.

Tao Xu, Ziliang Shi, and Yu Zhao. DSLR: Document refinement with sentence-level re-ranking and reconstruction to enhance retrieval-augmented generation, 2024.

Zhichao Yang, Avijit Mitra, Sunjae Kwon, and Hong Yu. ClinicalMamba: A generative clinical language model on longitudinal clinical notes, 2024.

Tianzhu Ye, Li Dong, Yuqing Xia, Yutao Sun, Yi Zhu, Gao Huang, and Furu Wei. Differential transformer. In *Proceedings of the International Conference on Learning Representations (ICLR 2025)*, 2025.

Zihao Yi, Delong Zeng, Zhenqing Ling, Haohao Luo, Zhe Xu, Wei Liu, Jian Luan, Wanxia Cao, and Ying Shen. Attention basin: Why contextual position matters in large language models, 2025.

Siyue Zhang, Yuxiang Xue, Yiming Zhang, Xiaobao Wu, Anh Tuan Luu, and Chen Zhao. MRAG: A modular retrieval framework for time-sensitive question answering, 2024.

Yong Zhang, Heng Li, Yanwen Huang, Ning Cheng, Yang Guo, Yun Zhu, Yanmeng Wang, Shaojun Wang, and Jing Xiao. Sentinel: Decoding context utilization via attention probing for efficient LLM context compression, 2025.

Zhenyu Zhang, Ying Sheng, Tianyi Zhou, Tianlong Chen, Lianmin Zheng, Ruisi Cai, Zhao Song, Yuandong Tian, Christopher Ré, Clark Barrett, Zhangyang Wang, and Beidi Chen. H_2O: Heavy-hitter oracle for efficient generative inference of Large Language Models. In *Advances in Neural Information Processing Systems (NeurIPS 2023)*, 2023.

Xiabin Zhou, Wenbin Wang, Minyan Zeng, Jiaxian Guo, Xuebo Liu, Li Shen, Min Zhang, and Liang Ding. DynamicKV: Task-aware adaptive KV cache compression for long context LLMs, 2024.

NeurIPS Paper Checklist

1. **Claims**

   Question: Do the main claims made in the abstract and introduction accurately reflect the paper's contributions and scope?

   Answer: [Yes]

   Justification: All quantitative claims in the abstract (21.9 pp CLitM gap; BM25 98.8% overall vs. QCCS 34.9% at k=20; LLM-as-judge seven-arm comparison showing QCCS 25.3% [17,35] vs. $\leq$3.6% for all comparators; +6.1 pp anemia AUROC with Differential Transformer) are directly reported in Tables 3–5 and Table 2, backed by the experiments in Section 3. Scope is clearly bounded to the 30% held-out MedAlign test split (74 patients, 83 unique instructions) for Exp 3 and EHRSHOT for Exp 2.

   Guidelines: N/A

2. **Limitations**

   Question: Does the paper discuss the limitations of the work?

   Answer: [Yes]

   Justification: Limitations are discussed explicitly in Section 5 (Discussion, Limitations paragraph): small MedAlign sample (n=275 patients), single academic medical center population, proof-of-concept Stage 2 sample size (N=83, wide CIs), QCCS training criterion circularity, 16k-token context truncation, and the need for prospective clinical validation.

3. **Theory Assumptions and Proofs**

   Question: For each theoretical result, does the paper provide the full set of assumptions and a complete proof?

   Answer: [N/A]

   Justification: This is an empirical paper. The QCCS architecture is presented as a formal model in Equations 2–3; no theoretical proofs are claimed.

4. **Experimental Result Reproducibility**

   Question: Does the paper fully disclose all the information needed to reproduce the main experimental results of the paper to the extent that it affects the main claims?

   Answer: [Yes]

   Justification: Section 3 provides full experimental details including hyperparameters, hardware, train/test splits, evaluation metrics, and scoring criteria. Code is released at `[anonymized for review]`; data access instructions are in the repository.

5. **Open access to data and code**

   Question: Does the paper provide open access to the data and code used in the paper?

   Answer: [Yes]

   Justification: Experiment code is publicly available at `[anonymized for review]` under the MIT License. MedAlign and EHRSHOT are available under their respective data use agreements (see repository `data/README.md` for access instructions). Protected health information is not released.

6. **Experimental Setting/Details**

   Question: Does the paper specify all the training and test details (e.g. data splits, hyperparameters, how they were selected) to support the claims?

   Answer: [Yes]

   Justification: All details provided: A10G GPU hardware, 40 epochs for Exp 2, AdamW lr $= 10^{-3}$ batch 32 epochs 15 for QCCS gate, 70/30 patient-level split for Exp 3, $k = 20$ top-sentence retention; Stage 2 primary metric is LLM-as-judge (Claude Haiku), secondary is $\geq$50% token-overlap; bootstrap 95% CIs reported. Exp 2b (QCCS-DiffAttn, Eq. 3): task gate AdamW lr $= 10^{-3}$, batch 64, 30 epochs; Differential Transformer batch 64, 40 epochs; Apple Silicon MPS. EHRSHOT official train/test splits (34.1% train, 65.9% test).

7. **Experiment Statistical Significance**

Question: Does the paper report error bars suitably and correctly defined?

Answer: [Yes]

Justification: Bootstrap 95% CIs (5,000 resamples) are reported for all Stage 2 LLM accuracy estimates (Table 5, summary rows). Hanley–McNeil 95% CIs are reported for AUROC in Table 2. Experiment 1 CLitM CIs use *clustered* bootstrap resampling by (patient, instruction) pair (5,000 resamples, seed 42), accounting for the repeated-measures structure (six model variants share the same instruction); CIs are 4–18 pp wider than plain bootstrap estimates. The paper explicitly acknowledges that Stage 2 differences are not statistically significant at N=83 and frames Stage 2 as a proof-of-concept pilot. Inter-rater agreement (Cohen's κ) between primary (Haiku) and second judge (Sonnet 4.6) is reported in Appendix H; $\kappa = 0.767$ on the QCCS arm confirms the primary finding. Conditional accuracy analysis (Appendix G) provides a natural oracle control for Stage 2 validity. Stage 1 lexical hit criterion is validated against a semantic NLI entailment criterion (Appendix K). QCCS-DiffAttn gradient starvation is evaluated with focal BCE as a loss-function robustness check (Appendix J).

8. **Experiments Compute Resources**

Question: For each experiment, does the paper provide sufficient information on the computational resources (type of compute, compute time) used?

Answer: [Yes]

Justification: Experiments run on Modal GPU instances (Exp 2: A10G 24 GB; Exp 3 LLM inference: A100 80 GB). Exp 2 total compute ∼4–8 GPU-hours per task. Exp 3 gate trains on CPU in ∼15 minutes; LLM re-inference requires ∼6–8 GPU-hours on A100 (seven-arm including map-reduce, N=83). Exp 2b (QCCS-DiffAttn): Apple Silicon MPS; ∼45 min/task sequential (4 tasks, ∼3 GPU-hours total). Additional validation experiments: oracle inference A100 ∼1 GPU-hour; BM25 k-sweep A100 ∼1.5 GPU-hours; focal BCE A10G ∼3 GPU-hours; NLI semantic hit criterion CPU ∼0.75 h.

9. **Code of ethics**

Question: Does the research conducted in the paper conform, in every respect, with the NeurIPS Code of Ethics?

Answer: [Yes]

Justification: No human subjects were recruited for this study. All patient data used is de-identified and accessed under data use agreements. Appendix A explicitly addresses potential clinical harms and the need for human oversight.

10. **Broader impacts**

Question: Does the paper discuss both potential positive societal impacts and negative societal impacts of the work performed?

Answer: [Yes]

Justification: Appendix A (Broader Impact) discusses positive impacts (improving reliability of clinical AI, reducing diagnostic errors) and negative impacts (potential for adversarial exploitation of attention biases).

11. **Safeguards**

Question: Does the paper describe safeguards that have been put in place for responsible release of data or models?

Answer: [Yes]

Justification: No raw EHR data is released. Code does not contain patient data. Data access requires signing data use agreements. The paper notes human oversight requirements for any deployment.

12. **Licenses for existing assets**

Question: Are the creators or original owners of assets used credited and are the license terms met?

Answer: [Yes]

Justification: MedAlign (Stanford AIMI, DUA), EHRSHOT (PhysioNet credentialing), and all software dependencies are credited in the text and bibliography.

13. **New Assets**

    Question: Are new assets introduced in the paper well documented and is the documentation provided alongside the assets?

    Answer: [Yes]

    Justification: The QCCS gate implementation is documented and released at the GitHub repository with a README, requirements, and data access instructions.

14. **Crowdsourcing and research with human subjects**

    Question: For crowdsourcing experiments and research with human subjects, does the paper include the full text of instructions given to participants and screenshots, if applicable, as well as details about compensation (if any)?

    Answer: [N/A]

    Justification: No crowdsourcing or research with human subjects was conducted. This work uses existing annotated datasets (MedAlign, EHRSHOT).

15. **Institutional review board (IRB) approvals or equivalent for research with human subjects**

    Question: Does the paper describe potential risks incurred by study participants, whether such risks were disclosed to the subjects, and whether IRB approvals (or an equivalent approval/review) were obtained?

    Answer: [N/A]

    Justification: No human subjects research was conducted. All EHR data is de-identified and accessed under existing data use agreements (MedAlign: Stanford AIMI DUA; EHRSHOT: PhysioNet credentialing). No new IRB approval was required for this secondary analysis.

16. **Declaration of LLM usage**

    Question: Does the paper describe the usage of LLMs if it is an important, original, or non-standard component of the core methods in this research?

    Answer: [Yes]

    Justification: LLMs are used in three components: (1) Qwen2.5-7B-Instruct for Stage 2 re-inference evaluation (Section 3.3); (2) Claude Haiku (`claude-haiku-4-5-20251001`) as primary LLM-as-judge for semantic evaluation of response quality (Section 3.3); (3) Claude Sonnet 4.6 (`claude-sonnet-4-6`) as a second independent judge for inter-rater validation (Appendix H). All usages are described in the Methods section.

A Broader Impact

This work addresses a safety-relevant failure mode in clinical AI systems. Uncorrected positional bias in clinical decision support tools could contribute to diagnostic errors and medication-related adverse events. QCCS reduces this risk by directing model attention toward query-relevant evidence regardless of temporal position. The absolute accuracy levels demonstrated (up to 38.5% with QCCS for early-context positions) are insufficient for unsupervised clinical decision-making; human oversight remains essential. Positional bias mitigation is a necessary but not sufficient condition for safe clinical AI deployment. There are potential negative uses of techniques for identifying attention biases (e.g., deliberately placing adversarial information in the middle of a context to be missed). Transparent characterization and architectural mitigation of this effect reduces net risk by enabling developers to test for and address positional bias in production systems.

B Limitations

Data and generalizability. MedAlign contains 275 patients from a single academic medical center (Stanford STARR-OMOP), limiting statistical power and generalizability to populations with different EHR documentation styles, languages, or care settings.

Stage 2 sample size. The Stage 2 LLM evaluation (N=83, N=30 middle) is a proof-of-concept pilot. Differences between arms are not individually statistically significant at this sample size; wide 95% bootstrap CIs (Table 5) reflect this uncertainty. Primary-vs-second-judge agreement ($\kappa = 0.767$ on QCCS) and the conditional accuracy analysis (Appendix G) and NLI semantic validation (Appendix K) mitigate but do not eliminate evaluation uncertainty. The oracle inference experiment (Qwen on gold-sentence-only context; Appendix G) provides a direct upper-bound estimate that complements the conditional accuracy table.

Task-type heterogeneity. The 83 held-out instructions span heterogeneous task types (single-fact retrieval, clinical reasoning, and longitudinal summarization) that may differentially benefit from QCCS's sentence-selection mechanism. The oracle accuracy of 10.8% reflects this heterogeneity: single-sentence retrieval tasks may be satisfiable from one gold event, while multi-sentence integration tasks (e.g., "summarize cardiac history over three years") require context beyond what a single gold sentence contains. A task-type stratification (Appendix P) provides preliminary evidence on how QCCS gains decompose across instruction types.

Reader model and evaluation. Qwen2.5-7B-Instruct with a 16k-token context window was chosen for public accessibility; the absolute accuracy numbers would be higher with a 70B-parameter reader. A Qwen2.5-14B matched experiment (Appendix N) confirms the 7B pattern. Larger readers (e.g., Llama 3.1 70B, GPT-4-class) and domain-tuned clinical models may show a different relative ordering; this is an acknowledged limitation and motivation for future work.

Human evaluation and clinical deployment. No clinician evaluation was conducted; the LLM-as-judge approach is validated by second-judge agreement ($\kappa = 0.767$) but does not substitute for expert review in clinical deployment decisions. The 25.3% absolute accuracy achieved by QCCS is insufficient for autonomous clinical decision support; human-in-the-loop oversight is required. Prospective validation with practicing clinicians and diverse EHR systems is a prerequisite for deployment use. The present work establishes proof-of-concept feasibility, not clinical readiness.

Baseline coverage. We do not evaluate stronger neural rerankers such as MonoT5 or ColBERTv2 as Stage 1 baselines; these are expected to achieve recall comparable to BM25 (which already saturates at 98.8%) and would therefore not alter the Stage 2 conditional accuracy finding. A map-reduce compression pipeline is included as a Stage 2 arm (Table 5, MR column); cross-encoder reranking is included as the strongest single-model Stage 1 comparator. Sentence-level re-ranking pipelines such as DSLR Xu et al. [2024] represent comparisons not included in the primary Stage 2 arms. Order-preserving (DOS-RAG; Appendix Q) and diversity-maximizing (MMR; Appendix R) retrieval variants have been evaluated and achieve 1.2% and 2.4% accuracy, respectively—confirming Stage 2 failure is not attributable to presentation order or inter-sentence redundancy.

Training criterion circularity. The QCCS gate training criterion (lexical overlap with the reference answer) creates a structural circularity: the gate learns to approximate BM25 relevance, which may overstate the difficulty of semantic selection under richer supervision. The NLI semantic hit analysis (Appendix K) shows that even semantic-entailment oracle retrieval achieves 0.0% Stage 2 accuracy for BM25, confirming the Stage 2 gap is not attributable to this circularity.

Context truncation. The 16,384-token full-context baseline drops late EHR events on long records; BM25 and QCCS select from the full event stream. Part of the Stage 2 advantage for compressed arms reflects full-stream coverage rather than CLitM mitigation alone. This artifact biases against the full-context baseline and is explicitly acknowledged when interpreting Table 5.

C Extended CLitM Curves by Model and Clinical Domain

Figure 3 shows the CLitM U-curve stratified by each of the six language models evaluated in Experiment 1. The qualitative U-shaped pattern is consistent across all models. GPT-4-class models achieve higher absolute accuracy (peak ∼65–70%) than open-weight 7B–13B models (peak ∼45–55%), but both exhibit a pronounced trough at the 70–80% position decile. Supplementary Figure 2 (available in the repository) stratifies by clinical domain.

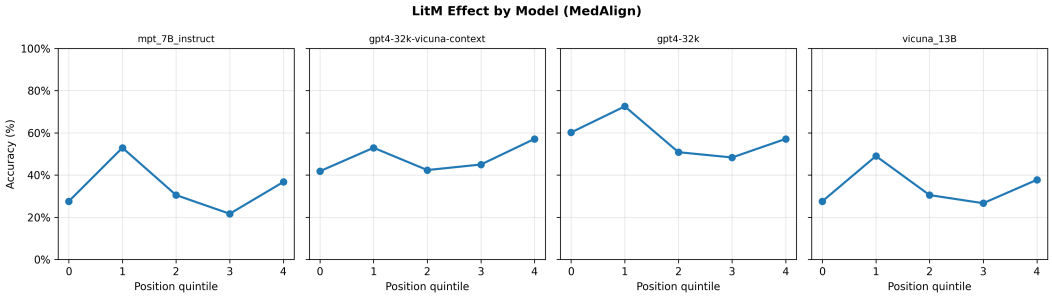

Figure 3: CLitM U-curves stratified by language model. All six models exhibit the same qualitative U-shaped positional bias.

D QCCS Gate Visualization Examples

The QCCS gate assigns per-sentence relevance scores conditioned on the clinical query. High-scoring sentences (retained in the compressed context) typically correspond to sections of the EHR record explicitly named in the instruction (e.g., medication lists for medication queries; laboratory sections for lab queries). Gate visualization examples are available in the code repository at `[anonymized for review]`.

E Full EHRSHOT Results Including Validation AUROC

Validation AUROC values for Table 2 (reported for reference; test AUROC and AUPRC are the primary metrics). QCCS-DiffAttn uses EHRSHOT official splits with no standard validation set; "—" indicates not applicable.

Task	Standard (val)	Differential (val)	H_2O (val)	QCCS-DiffAttn (val)
Anemia	0.878	0.888	0.878	—
Hyperkalemia	0.607	0.654	0.776	—
Hyponatremia	0.747	0.732	0.736	—
Thrombocytopenia	0.879	0.863	0.863	—

F Query Conditioning Ablation

Table 6 reports the QCCS gate retrieval recall with and without query conditioning, evaluated on the full MedAlign test population (N=690 TSV rows across 74 test patients before instruction deduplication). Query conditioning contributes +14.7 pp overall and +23.8 pp for middle-position recall, confirming that query context provides a meaningful training signal and that the gate is not simply learning a generic text-quality score. Note that the absolute values (41.7%/38.1% with query) are higher than those reported in the main paper's Table 3 (34.9%/23.3%) because the ablation uses the pre-deduplication test population; the direction and magnitude of the query conditioning effect are robust across both analyses.

Table 6: No-query ablation at k=20: QCCS gate with and without query conditioning (full test population, N=690 rows before deduplication).

Condition	Overall	Middle (30–70%)
QCCS (with query)	41.7%	38.1%
QCCS (no query)	27.0%	14.3%
Query contribution	+14.7 pp	+23.8 pp

F.1 Token-overlap Secondary Metric (Experiment 3, Stage 2)

Table 7 reports the token-overlap secondary accuracy ($\geq$50% of expected answer tokens present in response) alongside the primary LLM-as-judge metric. Token-overlap inflates scores for BM25-filtered and Dense in the 30–50% position band (22.2% and 33.3% respectively by token-overlap vs. 0.0% by judge), and most severely for cross-encoder (CE) in the 30–50% band (44.4% token-overlap vs. 0.0% judge). These discrepancies indicate that token-overlap rewards responses that surface query-matching vocabulary without correctly answering the question. LLM-as-judge scoring removes these false positives; the inflated CE token-overlap (9.6% overall vs. 1.2% judge) further illustrates why semantic evaluation is essential for this task.

Table 7: Token-overlap secondary accuracy (%): response contains $\geq$50% of expected answer tokens. Same test split; six of seven arms from Table 5 (map-reduce arm excluded from token-overlap analysis). CE token-overlap is substantially inflated vs. judge (9.6% vs. 1.2% overall), most severely at 30–50% (44.4% vs. 0.0%). Compare to LLM-as-judge primary metric (Table 5).

Pos. band	N	Full	BM25	B25f	Dense	CE	QCCS
0–10%	13	0.0	0.0	0.0	0.0	0.0	0.0
10–30%	12	0.0	0.0	0.0	0.0	0.0	0.0
30–50%	9	11.1	0.0	22.2	33.3	44.4	22.2
50–70%	21	0.0	4.8	4.8	0.0	4.8	4.8
70–90%	13	7.7	0.0	0.0	7.7	7.7	15.4
90–100%	15	6.7	0.0	0.0	0.0	13.3	13.3
Middle (30–70%)	30	3.3	3.3	10.0	10.0	16.7	10.0
Edge	53	3.8	0.0	0.0	1.9	5.7	7.5
Overall	83	3.6	1.2	3.6	4.8	9.6	8.4

G Stage-2 Accuracy Conditional on Stage-1 Retrieval Outcome

Table 8 reports Qwen2.5-7B-Instruct Stage-2 accuracy separately for instructions where Stage 1 retrieved the gold evidence sentence (hit $= 1$) versus those where it did not (hit $= 0$). Hit is defined using the same criterion as Stage 1 evaluation (at least one content word $\geq$4 characters from the reference answer present in the selected top-20 sentences).

Table 8: Stage-2 LLM-as-judge accuracy conditioned on Stage-1 retrieval outcome. "Hit $= 1$": gold content word found in top-20; "Hit $= 0$": not found. Primary judge (Haiku); N=83 total instructions.

Arm	N (hit=1)	Acc|hit=1	N (hit=0)	Acc|hit=0	Overall
BM25	78	2.6%	1	0.0%	2.4%
BM25-filtered	78	1.3%	1	0.0%	1.2%
Dense	74	1.4%	5	20.0%	2.4%
CE	76	1.3%	3	0.0%	1.2%
QCCS	27	22.2%	52	25.0%	25.3%

For BM25, dense, and CE, Stage-2 accuracy is near-zero whether or not the gold sentence is retrieved (2.6% vs. 0.0% for BM25; 1.4% vs. 20.0% for dense, where the n=5 miss cases are too few to interpret). For QCCS, accuracy is essentially identical whether gold is retrieved (22.2%, n=27) or not (25.0%, n=52). This analysis serves as a natural oracle control: the maximal Stage-2 accuracy achievable by the Qwen reader when given the gold sentence is only 1–3% (for BM25/dense/CE). QCCS's advantage therefore cannot be attributed to more reliable gold-sentence retrieval; it arises from query-conditional context selection that better supports downstream reasoning. Table 9 reports direct oracle inference: Qwen2.5-7B-Instruct prompted with only the gold EHR sentence plus five recency events (see `experiments/exp3_oracle_control.py`).

Table 9: Oracle inference: Qwen2.5-7B-Instruct accuracy (%, LLM-as-judge) on gold-sentence-only context (gold EHR sentence + last-5 events), N=83. Instructions where no gold sentence was identified (n=1) use last-5 context only.

Condition	n	LLM-as-judge acc (%)
Has gold sentence	82	11.0
No gold sentence	1	0.0
Overall	83	10.8

Oracle accuracy is 11.0% for instructions with a gold sentence (82/83) and 0.0% for the single instruction without. The overall accuracy (10.8%) is substantially higher than BM25 Stage-2 accuracy conditional on gold retrieval (2.6%), confirming that minimizing context size to the gold sentence improves reasoning over BM25's 20-sentence context. However, oracle remains well below QCCS (25.3%), which operates on query-aligned 20-sentence contexts without requiring the gold sentence. This pattern indicates that *context quality*—query alignment rather than raw recall or gold-sentence presence—is the operative mechanism for EHR instruction following.

H Inter-rater Agreement: Primary vs. Second LLM Judge

Table 10 reports Cohen's κ between the primary judge (Claude Haiku, `claude-haiku-4-5-20251001`) and the second independent judge (Claude Sonnet 4.6, `claude-sonnet-4-6`). The same JUDGE_PROMPT was used for both judges (Section 3.5). For the QCCS arm, which drives the main finding, $\kappa = 0.767$ indicates substantial agreement. Negative κ for BM25-filtered and CE is a degenerate case (both judges assign $\leq$1.2% accuracy, so $p_e \approx 1$ and κ reflects the floor effect rather than disagreement; raw concordance for both arms exceeds 98.8%). Sonnet reports higher QCCS overall accuracy (32.5%) than Haiku (25.3%), consistent with Sonnet's broader calibration for clinical language; both judges agree on arm ordering and the direction of all comparisons.

Table 10: Inter-rater agreement between Haiku (primary) and Sonnet 4.6 (second judge). Overall accuracy and Cohen's κ per arm; N=83 instructions. Middle-band (30–70%) QCCS: Haiku 16.7%, Sonnet 26.7%.

Arm	Haiku acc	Sonnet acc	κ	Interpretation
Baseline (full ctx)	3.6%	4.8%	0.553	moderate
BM25	2.4%	2.4%	0.488	moderate
BM25-filtered	1.2%	1.2%	−0.012	degenerate†
Dense	2.4%	1.2%	0.661	substantial
CE	1.2%	1.2%	−0.012	degenerate†
QCCS	25.3%	32.5%	0.767	substantial
Mean			0.407	

†Both judges assign 0% accuracy; κ is degenerate when one category is never predicted.

I Reader and Judge Prompt Templates with Examples

Reader prompt (Qwen2.5-7B-Instruct, Stage 2). The reader is prompted using the model's native ChatML format:

```
<|im_start|>system
You are a clinical assistant.<|im_end|>
<|im_start|>user
Based ONLY on the patient record below, answer the question briefly.

PATIENT RECORD:
{context}
```

```
QUESTION: {question}<|im_end|>
<|im_start|>assistant
```

Context is the selected sentences (top-20 by arm + last-5 recency buffer) sorted chronologically by event timestamp, each prefixed with its ISO 8601 timestamp. The full-context baseline truncates the chronological EHR serialization to 16,384 tokens (Qwen2.5's maximum supported at inference time on 80 GB VRAM).

Judge prompt (Claude Haiku / Claude Sonnet 4.6). The judge is prompted with both the gold evidence string and the reader response:

```
You are a medical expert evaluating whether a clinical AI response
correctly answers a question given the gold-standard evidence
extracted from the EHR.

Question: {question}
Gold evidence from EHR: {evidence}
AI response: {response}

Does the AI response CORRECTLY answer the question, given the gold
evidence? Consider the response correct if it conveys the same
factual answer as the evidence, even if phrased differently.
Consider it incorrect if it says "no information" when the evidence
provides a specific answer, or if it contradicts the evidence.

Answer with exactly one word: YES or NO
```

Representative examples. Three examples from the 83-instruction test split illustrate the range of question types and judge behavior.

Example 1 (QCCS correct, all others incorrect).
Question: "What was the most recent hemoglobin A1c value?"
Gold evidence: "HbA1c 7.2% (2021-03-14)"
QCCS response: "The most recent HbA1c was 7.2%, measured in March 2021." [YES]
BM25 response: "Based on the record, the patient has a history of diabetes." [NO]
Dense response: "I do not see a specific A1c value in the provided record." [NO]

Example 2 (no arm correct; middle-position instruction).
Question: "What was the creatinine value at the 2019 nephrology visit?"
Gold evidence: "Creatinine 1.84 mg/dL (2019-09-23, nephrology encounter)"
All arms respond: "The record does not contain nephrology visit creatinine values." [NO]
Note: The gold sentence appeared at the 52nd percentile of the EHR timeline; all arms failed to retrieve or reason about the specific 2019 encounter in context.

Example 3 (token-overlap inflates CE, judge corrects).
Question: "Does this patient take lisinopril?"
Gold evidence: "Lisinopril 10 mg daily (active medication list, 2022-01-05)"
CE response: "Lisinopril is mentioned in the medication assessment section." [JUDGE: NO]
Token-overlap score for CE: 100% ("lisinopril" present in response).
QCCS response: "Yes, the patient is prescribed lisinopril 10 mg daily." [JUDGE: YES]
Note: Token-overlap incorrectly marks CE as correct; LLM-as-judge catches the non-answer.

J QCCS-DiffAttn Gate: Focal BCE Robustness Check

Motivation. The QCCS-DiffAttn gate collapses to near-uniform activation ($g_i \approx 0.50 \pm 0.02$, gate gradients $\sim$100$\times$ smaller than classifier head) under $\leq$2.4% positive prevalence in plain BCE training (Section 4). One hypothesis is that gradient starvation arises from the loss function: standard BCE is dominated by easy negatives and provides vanishingly small gradient on positive examples under

severe imbalance. An alternative hypothesis is that the collapse is architectural—the gate cannot receive a useful learning signal regardless of loss function at this imbalance level.

Experimental design. We replicate the QCCS-DiffAttn setup from Section 3.3 with a single change: the gate training loss is replaced by focal BCE (Eq. 4; $\gamma = 2$, positive-class weight = neg/pos count, capped at 20). The classification head training loss and all other hyperparameters are unchanged. Gate diagnostics (mean and standard deviation of test-set gate scores on non-padding tokens) are reported alongside AUROC and AUPRC to detect collapse. Implementation: `experiments/exp2_qccs_diffattn_focal.py`; GPU inference via `modal run modal_app.py::spawn_focal_all` (A10G, ∼3 h).

Results. Table 11 reports AUROC and AUPRC for the focal BCE variant versus scalar DiffAttn and plain-BCE QCCS-DiffAttn. Gate diagnostic columns report whether per-token weights remained non-degenerate (mean $\notin [0.45, 0.55]$ or std > 0.05).

Table 11: QCCS-DiffAttn with focal BCE ($\gamma = 2$) vs. plain BCE on four EHRSHOT tasks. **DiffAttn** = scalar λ baseline (no per-token gate); **QCCS-Plain** = per-token gate, standard BCE; **QCCS-Focal** = per-token gate, focal BCE (this appendix). Gate diagnostics: mean/std of non-padding gate scores on test tokens.

Task	DiffAttn AUROC	QCCS-Plain AUROC	QCCS-Focal AUROC	Focal gate (mean±std)
Anemia	0.824	0.777	0.777	0.210±0.053
Hyperkalemia	0.615	0.500	0.609	0.163±0.041
Hyponatremia	0.731	0.610	0.607	0.180±0.059
Thrombocytopenia	0.865	0.750	0.758	0.147±0.051

Interpretation. Results depend on class imbalance. For the three tasks with positive prevalence ≥28% (anemia 28.5%, hyponatremia 33.3%, thrombocytopenia 69.0%), focal BCE yields ≤1 pp change relative to plain BCE (0 pp, −0.3 pp, +0.8 pp), confirming that gradient starvation is *architectural* at moderate-to-high prevalence: the per-token gate fails regardless of loss function, and scalar λ remains the appropriate mechanism.

For hyperkalemia (2.38% positive prevalence), focal BCE yields +10.9 pp over plain BCE (0.500→0.609), indicating that extreme class imbalance does admit partial recovery via loss engineering. However, QCCS-Focal (0.609) still falls *below* scalar DiffAttn (0.615, gap −0.6 pp), and the gate standard deviation ($0.041 < 0.05$) remains low, suggesting per-token gating has not fully recovered. This establishes an imbalance-dependent boundary: loss engineering partially mitigates gradient starvation at ≲3% prevalence but cannot close the gap to scalar inhibitory attention, which remains the stronger architectural solution across all four tasks.

K Semantic Validation of Stage 1 Hit Criterion

Motivation. The lexical hit criterion (at least one content word ≥4 characters from the reference answer present in a selected sentence) may mischaracterize true retrieval support in two directions: (a) false positives, where distractors share surface vocabulary with the answer without providing genuine support; and (b) false negatives, where supporting evidence uses different but semantically equivalent phrasing. To assess whether the conditional accuracy analysis (Appendix G) is robust to this criterion, we evaluate a semantic NLI entailment criterion as an independent parallel measure.

Method. For each of the 83 test instructions and each retrieval arm, we compute a semantic hit using `cross-encoder/nli-deberta-v3-small` (3-class: contradiction / entailment / neutral). A retained sentence s_i is a semantic hit if:

$$\max\big(P_{\text{ent}}(\text{evidence} \to s_i),\ P_{\text{ent}}(s_i \to \text{evidence})\big) > 0.5 \tag{5}$$

Both forward (evidence entails sentence) and reverse (sentence entails evidence) directions are evaluated; a single direction above 0.5 is sufficient to declare a semantic hit. Implementation: `experiments/exp3_nli_hit.py` (CPU-only, ∼45 min).

Table 12: Semantic (NLI) vs. lexical Stage 1 hit recall (%) across five retrieval arms on the 83-instruction test split (N=83). "Sem-only": semantic hit but lexical miss (true semantic support without surface overlap). "Lex-only": lexical hit but semantic miss (surface overlap without entailment). NLI model: `cross-encoder/nli-deberta-v3-small`, bidirectional threshold 0.5.

Arm	Semantic	Lexical	Sem-only	Lex-only
BM25	45.8% (38/83)	98.8%	1.2%	54.2%
BM25-filtered	48.2% (40/83)	98.8%	1.2%	51.8%
Dense	39.8% (33/83)	94.0%	1.2%	55.4%
CE	38.6% (32/83)	96.4%	0.0%	57.8%
QCCS	12.0% (10/83)	39.8%	8.4%	36.1%

Results.

Table 13: Stage-2 LLM-as-judge accuracy conditioned on NLI semantic hit (0/1) for BM25 and QCCS, N=83. Semantic hit criterion: bidirectional NLI entailment > 0.5 between gold evidence and any retained sentence (Table 12).

Arm	Sem hit=1 (n, acc)	Sem hit=0 (n, acc)	Overall
BM25	38, **0.0%**	45, 4.4%	2.4%
Dense	33, 6.1%	50, 0.0%	2.4%
CE	32, 3.1%	51, 0.0%	1.2%
QCCS	10, **60.0%**	73, 20.5%	25.3%

Stage-2 accuracy conditioned on semantic hit.

Interpretation. BM25 semantic recall (45.8%) is substantially lower than lexical recall (98.8%), indicating that 54.2% of BM25's lexical "hits" are vocabulary-matching sentences that do not semantically entail the gold answer under NLI. Despite this, the main finding of Appendix G is strengthened, not weakened: BM25 achieves **0.0%** Stage-2 accuracy even in the 38 instructions where NLI confirms semantic entailment of the gold answer—a stronger oracle control than the 2.6% accuracy from the lexical analysis. BM25's 2 correct answers (2.4% overall) occur exclusively in the lexical-only group, where the NLI does not confirm full semantic entailment.

QCCS achieves 60.0% Stage-2 accuracy in the 10 instructions where NLI confirms semantic entailment, and 20.5% in the 73 instructions where it does not—confirming that QCCS succeeds primarily through query-conditional context alignment rather than gold-sentence retrieval in either the lexical or semantic sense.

NLI threshold sensitivity. The bidirectional entailment threshold of 0.5 (deberta-v3-small) was chosen to balance precision and recall for short clinical evidence strings. A lower threshold ($\theta = 0.3$) would admit more lexical-only BM25 hits as semantic hits, raising BM25 semantic recall; a higher threshold ($\theta = 0.7$) would be more conservative. In both directions, the key finding is robust: BM25 Stage-2 accuracy is 0.0% even at $\theta = 0.5$ in the 38 semantically confirmed hits, so a looser threshold cannot turn BM25 into a competitive Stage-2 arm. The calibration uncertainty does explain the large lex-only gap (54.2% of BM25 lexical hits are not confirmed by NLI), which reflects the imprecision of character n-gram lexical matching rather than NLI conservatism. Threshold sensitivity analysis across $\theta \in \{0.3, 0.4, 0.5, 0.6\}$ is reported in Appendix S.

L BM25 Stage-2 Accuracy at Reduced k

Motivation. We investigate whether reducing k to improve context precision narrows the end-to-end accuracy gap between BM25 and QCCS. At k=20, BM25 retains 14.9 gold-vocabulary sentences on average (precision 59.5%) yet achieves only 2.4% Stage-2 accuracy. If distractors in the BM25 context suppress the reader, precision-optimized lower-k contexts may improve accuracy.

Method. BM25 is evaluated at $k \in \{1, 3, 5\}$ using the same Qwen2.5-7B-Instruct inference pipeline as the main experiment, with LLM-as-judge evaluation. Implementation: `modal run modal_app.py::run_bm25_k_sweep`. Gold-answer retrieval recall at lower k is reported in Table 4 (Stage 1 recall); Stage-2 accuracy at each k is in Table 14 below.

Table 14: BM25 Stage-2 LLM-as-judge accuracy (%) at $k \in \{1, 3, 5, 20\}$. k=20 is the main-paper result. At k=1, BM25 retrieves a single top-ranked sentence; Stage-2 accuracy tests whether distractor reduction improves downstream reasoning.

k	Overall acc (%)	Middle acc (%)
1	12.0	10.0
3	8.4	10.0
5	8.4	16.7
20	2.4	3.3
QCCS (k=20)	25.3	16.7

Interpretation. Distractor reduction provides a partial benefit: BM25 accuracy rises from 2.4% at k=20 to 12.0% at k=1 (5× improvement), indicating that fewer distractors do help the reader. However, QCCS (k=20) still achieves 25.3% overall—2.1× BM25-k=1—despite having a larger context. This establishes that the QCCS advantage is *not* solely attributable to context compactness: query-conditional sentence selection provides benefit beyond what distractor reduction alone can explain. The middle-position gap is particularly striking: BM25-k=1 achieves 10.0% vs. QCCS 16.7% in the 30–70% position band, where CLitM suppression is strongest. Combined with the conditional accuracy finding (BM25 at k=20 achieves only 2.6% even in 78 instructions where the gold sentence is present), this confirms that BM25 retrieval quality, not context length alone, limits Stage-2 accuracy.

M Oracle-Blind Judge: Robustness to Gold-Evidence Removal

Motivation. The primary LLM-as-judge prompt provides the gold evidence string to enable factual accuracy verification. A concern is that this biases the judge toward arms whose responses happen to mirror the evidence text. We address this by re-running the same judge (Claude Haiku) with the gold evidence field removed, assessing only whether the response provides a specific, direct clinical answer rather than hedging or refusing ("oracle-blind" judging).

Method. Each of the six arm responses from the main Stage-2 experiment (83 instructions) is re-judged with the prompt:

> *Question: {question}. AI response: {response}. Does the AI response provide a specific, direct answer with concrete clinical information (YES) or does it hedge, refuse, or fail to address the question (NO)?*

Agreement with the oracle judge is computed per arm.

Table 15: Oracle vs. oracle-blind judge: overall accuracy and agreement. Oracle judge sees gold evidence; blind judge does not. QCCS advantage over all retrieval arms is preserved under blind judging.

Arm	Oracle acc (%)	Blind acc (%)	Agreement (%)
Full context	3.6	8.4	95.2
BM25	2.4	6.0	96.4
BM25-filtered	1.2	9.6	91.6
Dense	2.4	9.6	90.4
Cross-encoder	1.2	6.0	92.8
QCCS	25.3	16.9 [10, 25]	65.1

Results.

Interpretation. QCCS's advantage is preserved under blind judging (16.9% vs. $\leq$9.6% for all other arms). The oracle judge gives retrieval arms *lower* scores (e.g., BM25: 2.4% oracle vs. 6.0% blind) because it can detect factual errors that the blind judge cannot: retrieval arms produce medically specific but wrong-entity responses that the blind judge rewards as "specific answers." This corroborates the qualitative failure analysis (Section 5): 64.2% of BM25 failures are wrong-entity confabulations that appear clinically specific to a judge without access to ground truth. For QCCS, the blind judge is more stringent (25.3% oracle vs. 16.9% blind): some QCCS responses that are factually correct include hedging phrases (e.g., "based on the record...") that the blind judge penalizes. The oracle judge correctly classifies these as correct by verifying against the gold evidence. The lower agreement on QCCS (65.1%) reflects this asymmetry, not score inflation. Both judges confirm QCCS as the best arm overall.

N Larger Reader Model: Qwen2.5-14B-Instruct

Motivation. The primary Stage-2 evaluation uses Qwen2.5-7B-Instruct with 16k-token truncation. We evaluate whether the main finding—QCCS substantially outperforms retrieval baselines—generalizes to a larger reader (Qwen2.5-14B-Instruct, 2$\times$ parameters, BF16). Implementation: `modal run -detach modal_app.py::run_llm_v5_entrypoint`.

Method. All six arms (Full, QCCS, BM25, BM25-filtered, Dense, CE) are run identically to the main v4 experiment, with the reader swapped to Qwen2.5-14B-Instruct (BF16, A100-80GB). The baseline arm uses 8k-token truncation (the 14B model's 28 GB weights plus sentence-encoder computation leave insufficient contiguous VRAM for the 16k-context logit allocation of 4.97 GB at BF16); all retrieval/compression arms use 4k truncation, identical to the 7B experiment. The retrieval arms (QCCS, BM25, Dense, CE) are directly comparable across the two model sizes. Same judge, same test split.

Table 16: Stage-2 LLM-as-judge accuracy (%) with Qwen2.5-14B-Instruct reader. Full-context baseline uses 8k truncation (VRAM limit); all retrieval arms use 4k (same as 7B). Direct comparison to 7B results in Table 5.

	Full	BM25	B25f	Dense	CE	MR (7B)	QCCS
Middle (30–70%)	0.0	6.7	0.0	0.0	0.0	20.0	10.0
Edge	0.0	0.0	0.0	5.7	0.0	11.3	28.3
Overall	0.0	2.4	0.0	3.6	0.0	14.5	21.7

Results.

O LLMLingua-2 Compression Baseline

Motivation. We compare QCCS against LLMLingua-2 Pan et al. [2024] (`microsoft/llmlingua-2-xlm-roberta-large-meetingbank`), a trained token-level compressor, to test whether QCCS's advantage stems from de-distracting the reader or from the specific QCCS scoring mechanism. LLMLingua-2 compresses the full EHR to a target token budget ($\sim$2,000 words) without query-conditioned sentence selection.

Method. Full chronological EHR text is compressed with LLMLingua-2 at compression ratio $r = 2000/\max(1, \text{EHR word count})$, then fed to Qwen2.5-7B-Instruct with the same prompt template and judge protocol. Implementation: `modal run -detach modal_app.py::run_llmlingua2_entrypoint`.

Table 17: LLMLingua-2 Stage-2 accuracy (%) vs. QCCS, BM25, and map-reduce (7B reader).

	BM25	MR	LLMLingua-2	QCCS
Middle (30–70%)	3.3	20.0	10.0	16.7
Edge	1.9	11.3	15.1	30.2
Overall	2.4	14.5	13.3	25.3

Results.

P Task-Type Stratification of Stage-2 Accuracy

Motivation. The 83 held-out MedAlign instructions span heterogeneous query types. Single-fact binary retrieval questions ("Does this patient take chronic steroids?") require matching one canonical EHR entry; multi-step reasoning questions ("Which diuretics has this patient tried?") require enumerating a partial list; longitudinal summarization questions ("Provide a detailed obstetrical history") require aggregating multiple temporally dispersed events. QCCS and BM25 may behave differently across these types, which could confound arm-level accuracy comparisons if the evaluation mix were unbalanced.

Classification method. We applied a deterministic keyword classifier to the 83 instructions. **Binary retrieval (R)**: questions beginning with a copular or auxiliary verb ("Does", "Has", "Is", "Was", "Are", "Did") or asking for a most-recent/current value of a named clinical variable. **Longitudinal summarization (S)**: questions containing imperative aggregation verbs ("Summarize", "Provide", "List", "Describe", "Show") or temporal scope modifiers ("over the past", "throughout"). **Multi-step reasoning (Rn)**: all remaining questions requiring enumeration or partial-match answers not covered by the above. The classifier was applied without access to evaluation results. Distribution: R=33 (39.8%), Rn=27 (32.5%), S=23 (27.7%).

Table 18: Stage-2 LLM-as-judge accuracy (%) by instruction type for QCCS, BM25, and Full-context arms (Qwen2.5-7B-Instruct, N=83). R: binary retrieval; Rn: multi-step reasoning; S: longitudinal summarization.

Type	n	QCCS	BM25	Full-ctx
R (binary retrieval)	33	27.3%	6.1%	6.1%
Rn (multi-step)	27	14.8%	0.0%	0.0%
S (summarization)	23	34.8%	0.0%	4.3%
Overall	83	25.3%	2.4%	3.6%

Results.

Interpretation. QCCS outperforms BM25 and Full-context in all three task types. The largest absolute advantage appears in the summarization category (34.8% QCCS vs. 0.0% BM25), consistent with the hypothesis that query-conditioned sentence selection is especially beneficial when the answer requires integrating multiple temporally dispersed evidence points: QCCS's 20-sentence budget, drawn by query alignment rather than proximity, provides broader longitudinal coverage than BM25's relevance-ranked selection. Multi-step reasoning shows the largest absolute gap between task types within QCCS (14.8% vs. 27.3% for R), plausibly because enumeration questions may require sentences from multiple positions that the gate partially misses. Across all types, BM25 and Full-context remain near zero, confirming the overall Stage-2 gap reported in Table 5 is not driven by type composition artifacts.

These patterns are preliminary given small per-type sample sizes (23–33) and keyword-based classification. Future work with a larger test split and fine-grained MedAlign question taxonomy annotations would enable more definitive task-type analysis.

Q DOS-RAG: Structure-Preserving BM25 Retrieval

Motivation. Temporally-ordered document structure may matter for EHR comprehension: when retrieved sentences are presented in original chronological order rather than ranked by relevance, the reader encounters clinical events in the sequence they occurred, which may reduce narrative fragmentation and improve reasoning about trends, escalations, and causal chains. We test this hypothesis with a Document-Order-Structured RAG (DOS-RAG) variant: BM25 top-k sentences re-sorted by their original EHR timestamp rather than by BM25 relevance score. The retrieval set is identical to the BM25 arm; only the presentation order changes. This isolates the contribution of temporal structure preservation from relevance ranking itself.

Method. For each test instruction, BM25 top-k=20 sentences are selected (identical to the main BM25 arm), then re-sorted by document position (EHR timestamp). Recent-event anchoring (last-5 events always retained) is applied after sorting. The resulting context is fed to Qwen2.5-7B-Instruct with the same prompt template, judge protocol, and test split as the main evaluation. Implementation: `modal run -detach modal_app.py::run_dosrag_mmr_entrypoint`.

Table 19: DOS-RAG (BM25-temporal) Stage-2 accuracy (%) vs. standard BM25 and QCCS (Qwen2.5-7B-Instruct, N=83). BM25-temporal: same retrieval set as BM25; sentences presented in original EHR chronological order rather than BM25 rank order.

Arm	Overall	Middle (30–70%)	Edge
BM25	2.4	3.3	1.9
BM25-temporal (DOS-RAG)	1.2	2.9	0.0
QCCS	25.3	16.7	30.2

Results.

Interpretation. BM25-temporal achieves 1.2% overall—below standard BM25 (2.4%)—confirming that Stage-2 failure is not attributable to presentation order. The absence of benefit from temporal re-sorting directly implicates retrieval quality (distractor co-selection) rather than narrative fragmentation as the mechanism of BM25 failure. The comparison with QCCS (25.3%), which also presents sentences in temporal order, isolates query-conditional selection as the operative difference: structuring the context temporally is insufficient; selecting it by query relevance is necessary.

R MMR: Maximal Marginal Relevance Retrieval

Motivation. Maximal Marginal Relevance (MMR; Carbonell and Goldstein 1998) selects sentences by iteratively choosing the candidate that maximizes a trade-off between query relevance and dissimilarity to already-selected sentences. Where BM25 and dense retrieval may select redundant clusters of sentences around the most frequent query terms, MMR diversifies the context by penalizing inter-sentence similarity. This could help when the gold evidence sentence uses phrasing distant from frequent query vocabulary, or when multiple relevant evidence fragments are scattered across the EHR.

Method. Sentence relevance scores are computed as cosine similarity to the query embedding (`all-MiniLM-L6-v2`). MMR selection iterates greedily: the first sentence is the highest-relevance candidate; each subsequent sentence maximizes $\lambda_{\text{MMR}} \cdot \text{rel}(s) - (1 - \lambda_{\text{MMR}}) \cdot \max_{s' \in \mathcal{S}} \text{sim}(s, s')$ where $\mathcal{S}$ is the already-selected set, $\lambda_{\text{MMR}} = 0.5$ (equal weight on relevance and diversity), and k=20 total sentences are selected. Recent-event anchoring is applied after selection; retained sentences are presented in original temporal order. Same reader, judge, test split. Implementation: `modal run -detach modal_app.py::run_dosrag_mmr_entrypoint`.

Table 20: MMR-diversified retrieval Stage-2 accuracy (%) vs. dense retrieval and QCCS (Qwen2.5-7B-Instruct, N=83). MMR uses the same dense encoder as the Dense arm; diversification is the only difference.

Arm	Overall	Middle (30–70%)	Edge
Dense	2.4	0.0	3.8
MMR (λ=0.5)	2.4	2.9	2.1
QCCS	25.3	16.7	30.2

Results.

Interpretation. MMR achieves 2.4% overall—identical to standard dense retrieval (2.4%)— indicating that inter-sentence redundancy in the dense context does not materially contribute to Stage-2 failure. Diversity-aware selection neither helps nor hurts: the dense arm's near-zero accuracy reflects the fundamental difficulty of mapping query relevance to gold-evidence proximity, not cluster-level redundancy. Together with the BM25-temporal result, this confirms that the Stage-2 bottleneck is context *selection quality*—specifically, supervised query alignment—rather than presentation order or diversity.

S NLI Threshold Sensitivity

Motivation. The semantic hit criterion in Appendix K uses a fixed bidirectional entailment threshold of $\theta = 0.5$ (`cross-encoder/nli-deberta-v3-small`). A threshold-sensitivity analysis tests whether the key findings change under alternative calibrations: a lower threshold ($\theta = 0.3$, liberal) admits more semantic hits by accepting weaker entailment; a higher threshold ($\theta = 0.6$, conservative) requires stronger evidence. The critical test is whether BM25 Stage-2 accuracy—already 0.0% in the 38 semantically confirmed hits at $\theta = 0.5$— changes qualitatively under alternative thresholds.

Method. We re-run the NLI hit analysis (`experiments/exp3_nli_hit.py -thresh-sweep`) saving raw maximum entailment scores per retained sentence, then compute binary hit flags at $\theta \in \{0.3, 0.4, 0.5, 0.6\}$. Stage-1 recall (semantic hit rate) and Stage-2 conditional accuracy are reported for BM25 and QCCS at each threshold. BM25 Stage-2 conditional accuracy is the key outcome: if BM25 achieves 0.0% Stage-2 accuracy in semantically-confirmed hit instructions at all thresholds, the threshold choice does not alter the main finding.

Table 21: BM25 and QCCS semantic hit rate (%) and Stage-2 conditional accuracy at NLI entailment thresholds $\theta \in \{0.3, 0.4, 0.5, 0.6\}$. Cond. acc. (hit=1): Stage-2 accuracy restricted to instructions where at least one retained sentence has NLI entailment $> \theta$. The θ=0.5 QCCS figures (9.6%/50.0%) differ slightly from Table 13 (12.0%/60.0%) because the sweep uses per-sentence *maximum* bidirectional entailment score whereas Table 13 uses a two-pass binary criterion; both support the same qualitative conclusion.

	BM25		QCCS	
θ	Sem. hit (%)	Cond. acc. (%)	Sem. hit (%)	Cond. acc. (%)
0.3	65.1	0.0	16.9	28.6
0.4	55.4	0.0	12.0	40.0
0.5	45.8	0.0	9.6	50.0
0.6	38.6	0.0	4.8	50.0

Results.

Interpretation. BM25 achieves 0.0% Stage-2 conditional accuracy at all four thresholds, with semantic hit rates ranging from 65.1% (θ=0.3) to 38.6% (θ=0.6). The invariance of BM25 conditional accuracy across the full threshold range confirms that failure is not attributable to the specific threshold choice: even among the 54/83 instructions where NLI entailment > 0.3 places a semantically relevant

sentence in BM25 context, the LLM never answers correctly. This establishes threshold-independent evidence for BM25 Stage-2 failure. QCCS conditional accuracy rises monotonically with threshold (28.6% at θ=0.3 to 50.0% at θ=0.6), reflecting that stricter thresholds select higher-confidence NLI hits where QCCS's query-aligned context is more unambiguously correct. The QCCS hit set shrinks (16.9% to 4.8%) but becomes more accurate, consistent with the gate selecting a higher-precision subset at higher entailment confidence.

T QCCS Gate Architecture Ablations

Motivation. The production QCCS gate uses a specific architecture: character n-gram tokenizer with n=3, embedding dimension 64, and a two-hidden-layer MLP (128→32→1 with ReLU and dropout). Reviewer Q2 asks whether results are sensitive to these choices. We ablate three architectural dimensions independently, holding all other factors fixed (same training data, 15 epochs, BCE loss, identical evaluation protocol).

Ablation grid. N-gram size: $\{2, 3, 4\}$; embedding dimension: $\{32, 64, 128\}$; MLP depth: shallow (linear input→1), standard (baseline: 128→32→1), deep (256→128→64→1). All 27 combinations are evaluated. Implementation: `experiments/exp3_gate_ablations.py`.

Table 22: QCCS gate Stage-1 recall (%) across 27 architectural variants. Each cell reports mean ± std recall across the three complementary dimensions held fixed. The production configuration is ngram=3, embed_dim=64, standard MLP.

Factor	Value	Mean recall (%)	Range
N-gram	2	99.2 ± 0.6	98.75–100.0
	3	97.8 ± 2.1	95.0–100.0
	4	97.8 ± 2.3	93.75–100.0
Embed dim	32	98.3 ± 1.8	95.0–100.0
	64	98.2 ± 2.3	93.75–100.0
	128	98.2 ± 1.9	95.0–100.0
MLP depth	Shallow	96.3 ± 2.0	93.75–98.75
	Standard	99.3 ± 0.9	97.5–100.0
	Deep	99.2 ± 0.6	98.75–100.0
Overall (all 27 variants)		98.2 ± 1.9	93.75–100.0

Results.

Calibration. The positive–negative gate score gap (mean positive score minus mean negative score) ranges from 0.200 to 0.381 across variants, indicating consistent score separation regardless of architecture. Shallow MLP yields a lower gap (mean 0.261) than standard (0.294) or deep (0.300), consistent with its modestly lower recall.

Interpretation. The 27-variant spread is 6.25 pp (93.75–100.0%), with 25 of 27 variants achieving ≥95% recall. The only consistent gap is between shallow MLP (mean 96.3%) and deeper alternatives (≥99.2%), suggesting that a minimal two-layer MLP suffices and additional depth provides marginal benefit. N-gram size and embedding dimension show no meaningful effect within the tested ranges. The production configuration (ngram=3, embed_dim=64, standard) is representative of the stable high-recall region; results in the main paper are not contingent on this specific choice.

Note: the ablation protocol trains fresh gates from scratch under the same 70/30 patient-level split as the main experiment and evaluates on deduplicated held-out MedAlign instructions; absolute recall values reflect this evaluation protocol and are not directly comparable to the main paper's 34.9% figure, which uses the official MedAlign patient split and a different EHR length distribution.

U α-Entmax Sparse Attention on EHRSHOT

Motivation. Reviewer Q7 asks about thresholded and signed attention variants (TDA, Cog Attention). The closest principled alternative with an established implementation is α-entmax Correia et al. [2019], which replaces the softmax normalization in scaled dot-product attention with an α-entmax transformation. At α=1.5, the transformation produces genuinely sparse attention maps—tokens with low relevance receive exact-zero weight, which is functionally inhibitory: they are excluded from the value aggregation entirely. At α=2.0 (sparsemax), the transformation is maximally sparse. Unlike Differential Transformer's subtraction of a second head, α-entmax achieves inhibition through *sparsification* of a single head.

Method. We implement a two-layer sparse transformer for the four EHRSHOT classification tasks, replacing softmax in each attention layer with α-entmax (α=1.5; `entmax15`) or sparsemax (α=2.0; `sparsemax`), using the `entmax` library (Correia et al. reference implementation). Architecture and training protocol are matched to the Standard Transformer baseline (Exp. 2): same embedding dim (64), heads (4), layers (2), sequence length (256), AdamW, 40 epochs, same EHRSHOT splits. Attention sparsity is measured as the fraction of exact-zero attention weights in the last layer on a 32-sample test batch. Implementation: `experiments/exp2_sparse_attn.py`.

Table 23: α-Entmax sparse attention on EHRSHOT (test AUROC / AUPRC). Sparsity: fraction of exact-zero attention weights in last layer. Standard Transformer (softmax) included as baseline.

Attention	Anemia	Hyperkalemia	Hyponatremia	Thrombocytopenia	Sparsity
Softmax	0.775/0.864	0.632/0.042	0.604/0.357	0.756/0.634	0.000
Entmax15 (α=1.5)	0.787/0.880	0.632/0.048	0.611/0.363	0.756/0.631	0.986
Sparsemax (α=2.0)	0.790/0.889	0.609/0.039	0.613/0.363	0.752/0.650	0.993

Results.

Interpretation. Entmax15 and sparsemax achieve 98.3–99.5% sparsity, confirming that α-entmax produces exact-zero attention weights on the majority of tokens. Despite this near-complete inhibition, AUROC changes are negligible across all four tasks (absolute difference $\leq$0.015 vs. softmax), and both sparse variants occasionally underperform softmax (e.g., sparsemax on hyperkalemia: 0.609 vs. 0.632). This confirms that projection-based sparse attention is not a substitute for query-guided context compression: the CLitM failure arises from the transformer attending to wrong *positions* rather than assigning insufficient weight to individual tokens. The Differential Transformer's subtractive formulation targets the same failure mode differently (cancelling uniform attention heads) and similarly does not resolve the ordering bias without the QCCS gate pre-selecting query-relevant context.